

W

Convolutional Networks (ConvNets): Part II

May 4, 2021

ddebarr@uw.edu

[http://cross-entropy.net/ml530/Deep Learning 3.pdf](http://cross-entropy.net/ml530/Deep_Learning_3.pdf)



https://en.wikipedia.org/wiki/Star_Wars_Day

Agenda for Tonight

- Homework Review
- [DLP] Chapter 5: Deep Learning for Computer Vision

The Three General Types of Network Layers

- Densely Connected Layers: each output value is based on all values from the previous layer (“global” features)

* We are here *

- Convolutional Layers: each output value is based on a spatially adjacent subset of values from the previous layer (“local” features)
- Embeddings and Recurrent Layers: soon [also Transformers]

Convolutional Operation Example

- Input Image: the 5x5, light-blue matrix
- Convolution Filter: the 3x3, dark-blue matrix
- Output Feature Map: the 3x3, green matrix
- The convolution filter is moved from left to right, and top to bottom, where each value in the output feature map is the sum of products between a 3x3 subimage of the input and the 3x3 convolution filter

3 ₀	3 ₁	2 ₂	1	0
0 ₂	0 ₂	1 ₀	3	1
3 ₀	1 ₁	2 ₂	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

Computing the 9 Output Values for the Filter

One output channel (matrix) per filter. The stride determines the step size when moving to the right or down

3 ₀	3 ₁	2 ₂	1	0
0 ₂	0 ₂	1 ₀	3	1
3 ₀	1 ₁	2 ₂	2	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3 ₀	2 ₁	1 ₂	0
0	0 ₂	1 ₂	3 ₀	1
3	1 ₀	2 ₁	2 ₂	3
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2 ₀	1 ₁	0 ₂
0	0	1 ₂	3 ₂	1 ₀
3	1	2 ₀	2 ₁	3 ₂
2	0	0	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0 ₀	0 ₁	1 ₂	3	1
3 ₂	1 ₂	2 ₀	2	3
2 ₀	0 ₁	0 ₂	2	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0 ₀	1 ₁	3 ₂	1
3	1 ₂	2 ₂	2 ₀	3
2	0 ₀	0 ₁	2 ₂	2
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1 ₀	3 ₁	1 ₂
3	1	2 ₂	2 ₂	3 ₀
2	0	0 ₀	2 ₁	2 ₂
2	0	0	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1	3	1
3 ₀	1 ₁	2 ₂	2	3
2 ₂	0 ₂	0 ₀	2	2
2 ₀	0 ₁	0 ₂	0	1

12	12	17
10	17	19
9	6	14

3	3	2	1	0
0	0	1	3	1
3	1 ₀	2 ₁	2 ₂	3
2	0 ₂	0 ₂	2 ₀	2
2	0 ₀	0 ₁	0 ₂	1

12	12	17
10	17	19
9	6	14

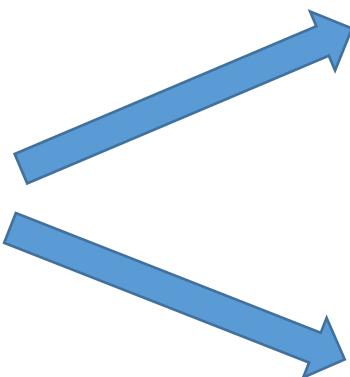
3	3	2	1	0
0	0	1	3	1
3	1 ₂	2 ₀	2 ₁	3 ₂
2	0 ₂	0 ₂	2 ₂	2 ₀
2	0 ₀	0 ₁	0 ₂	1 ₂

12	12	17
10	17	19
9	6	14

Pooling

pool_size = (2, 2) with strides = (2, 2) reduces (4, 4) “image” to (2, 2)

0.03	0.16	0.13	0.04
0.10	0.09	0.14	0.12
0.05	0.11	0.06	0.02
0.01	0.07	0.15	0.08



0.16	0.14
0.11	0.15

Max pooling

0.0950	0.1075
0.0600	0.0775

Average pooling

[DLP] Deep Learning for Computer Vision

1. Introduction to ConvNets
2. Training a ConvNet from Scratch on a Small DataSet
3. Using a PreTrained ConvNet
4. Visualizing What ConvNets Learn

Instantiation a Small ConvNet

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

>>> model.summary()
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
maxpooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
maxpooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
<hr/>		
Total params: 55,744		
Trainable params: 55,744		
Non-trainable params: 0		

Adding a Classifier on Top of the ConvNet

```
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

>>> model.summary()

Layer (type)                  Output Shape       Param #
=====                      =====
conv2d_1 (Conv2D)             (None, 26, 26, 32)    320
=====
maxpooling2d_1 (MaxPooling2D) (None, 13, 13, 32)    0
=====
conv2d_2 (Conv2D)             (None, 11, 11, 64)   18496
=====
maxpooling2d_2 (MaxPooling2D) (None, 5, 5, 64)      0
=====
conv2d_3 (Conv2D)             (None, 3, 3, 64)    36928
=====
flatten_1 (Flatten)           (None, 576)          0
=====
dense_1 (Dense)               (None, 64)           36928
=====
dense_2 (Dense)               (None, 10)           650
=====
Total params: 93,322
Trainable params: 93,322
Non-trainable params: 0
```

Training the ConvNet on MNIST Images

```
from keras.datasets import mnist
from keras.utils import to_categorical

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255

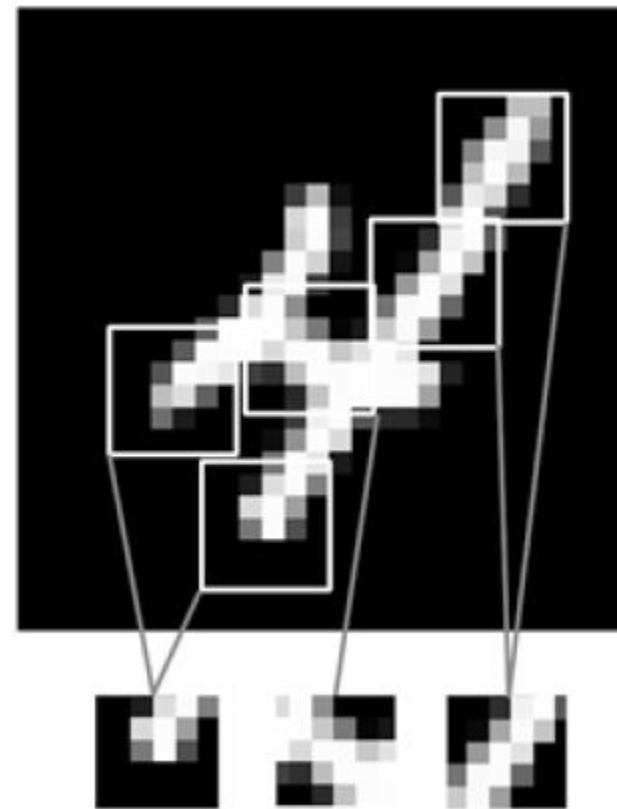
test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(train_images, train_labels, epochs=5, batch_size=64)
>>> test_loss, test_acc = model.evaluate(test_images, test_labels)
>>> test_acc
0.9908000000000001
```

Images Can Be Broken (Decomposed) Into Local Patterns

- Edges, Textures, etc.
- 5x5 patches

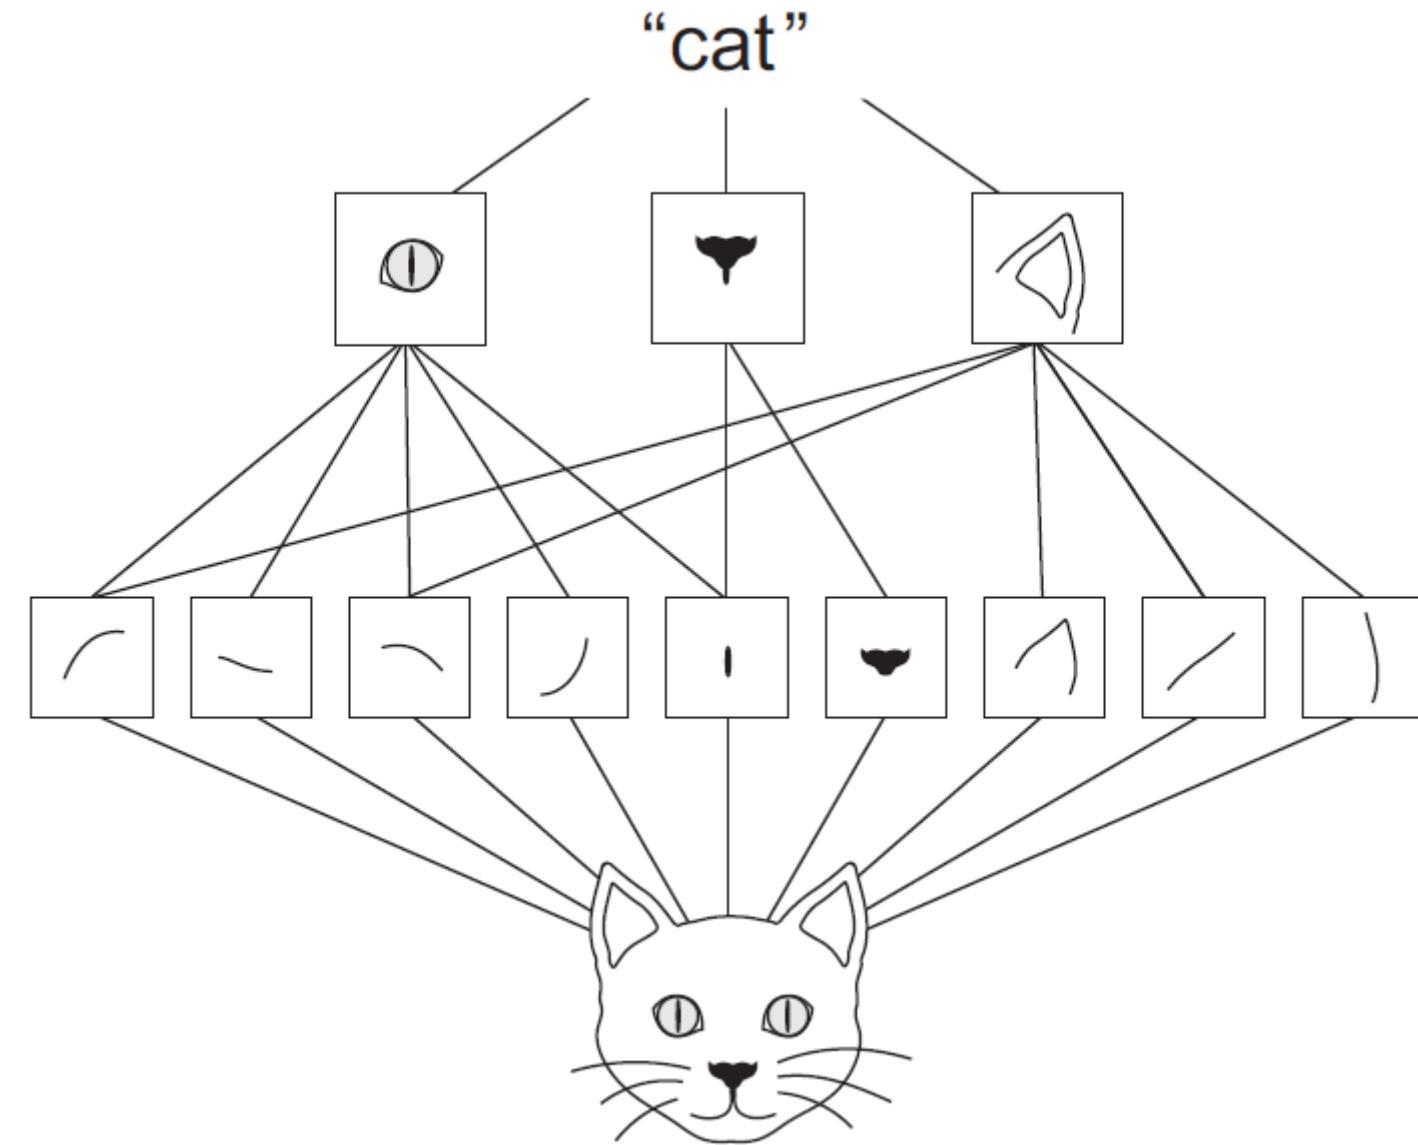


Local Patterns

- *The patterns they learn are translation invariant.* After learning a certain pattern in the lower-right corner of a picture, a convnet can recognize it anywhere: for example, in the upper-left corner. A densely connected network would have to learn the pattern anew if it appeared at a new location. This makes convnets data efficient when processing images (because *the visual world is fundamentally translation invariant*): they need fewer training samples to learn representations that have generalization power.
- *They can learn spatial hierarchies of patterns (see figure 5.2).* A first convolution layer will learn small local patterns such as edges, a second convolution layer will learn larger patterns made of the features of the first layers, and so on. This allows convnets to efficiently learn increasingly complex and abstract visual concepts (because *the visual world is fundamentally spatially hierarchical*).

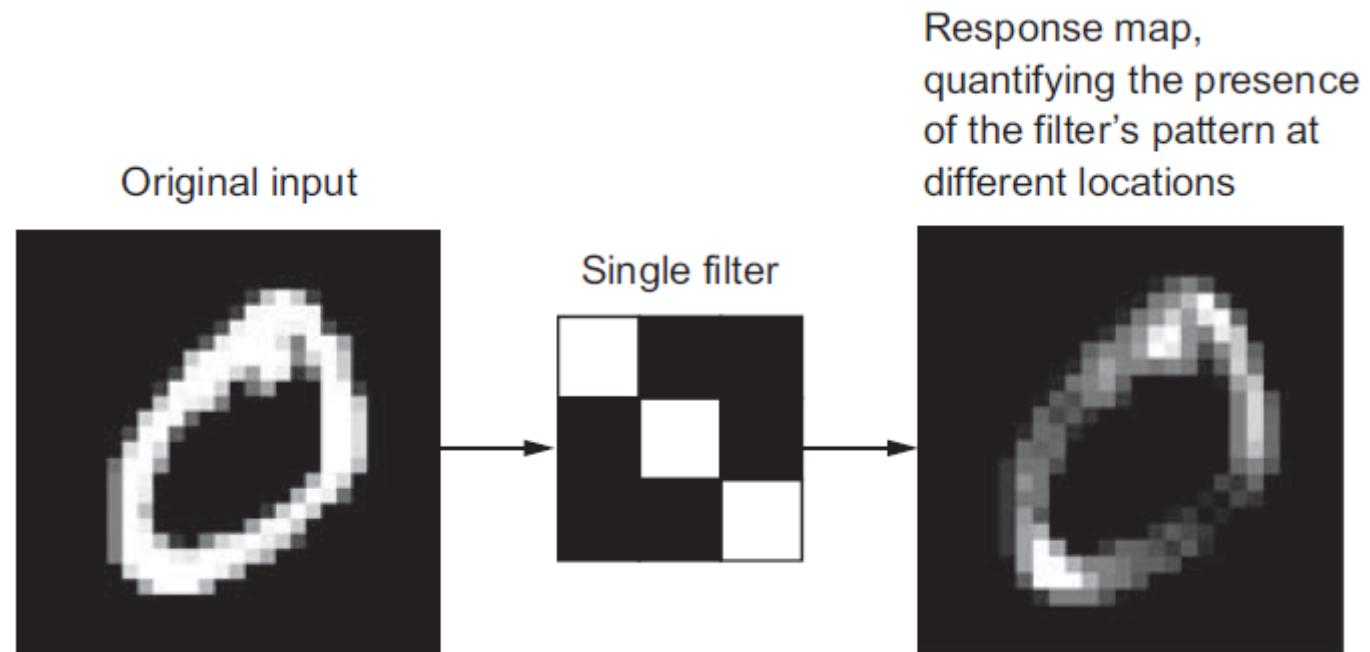
Local Edges Combine Into Local Objects

- Edges combine into objects, such as eyes, nose, and ears
- Eyes, nose, and ears combine into higher-level concepts, such as “cat”
- Read the picture from the bottom up [input to output]



Response Map

The Response Map is a two-dimensional map of the presence of a pattern at different feature locations in an input

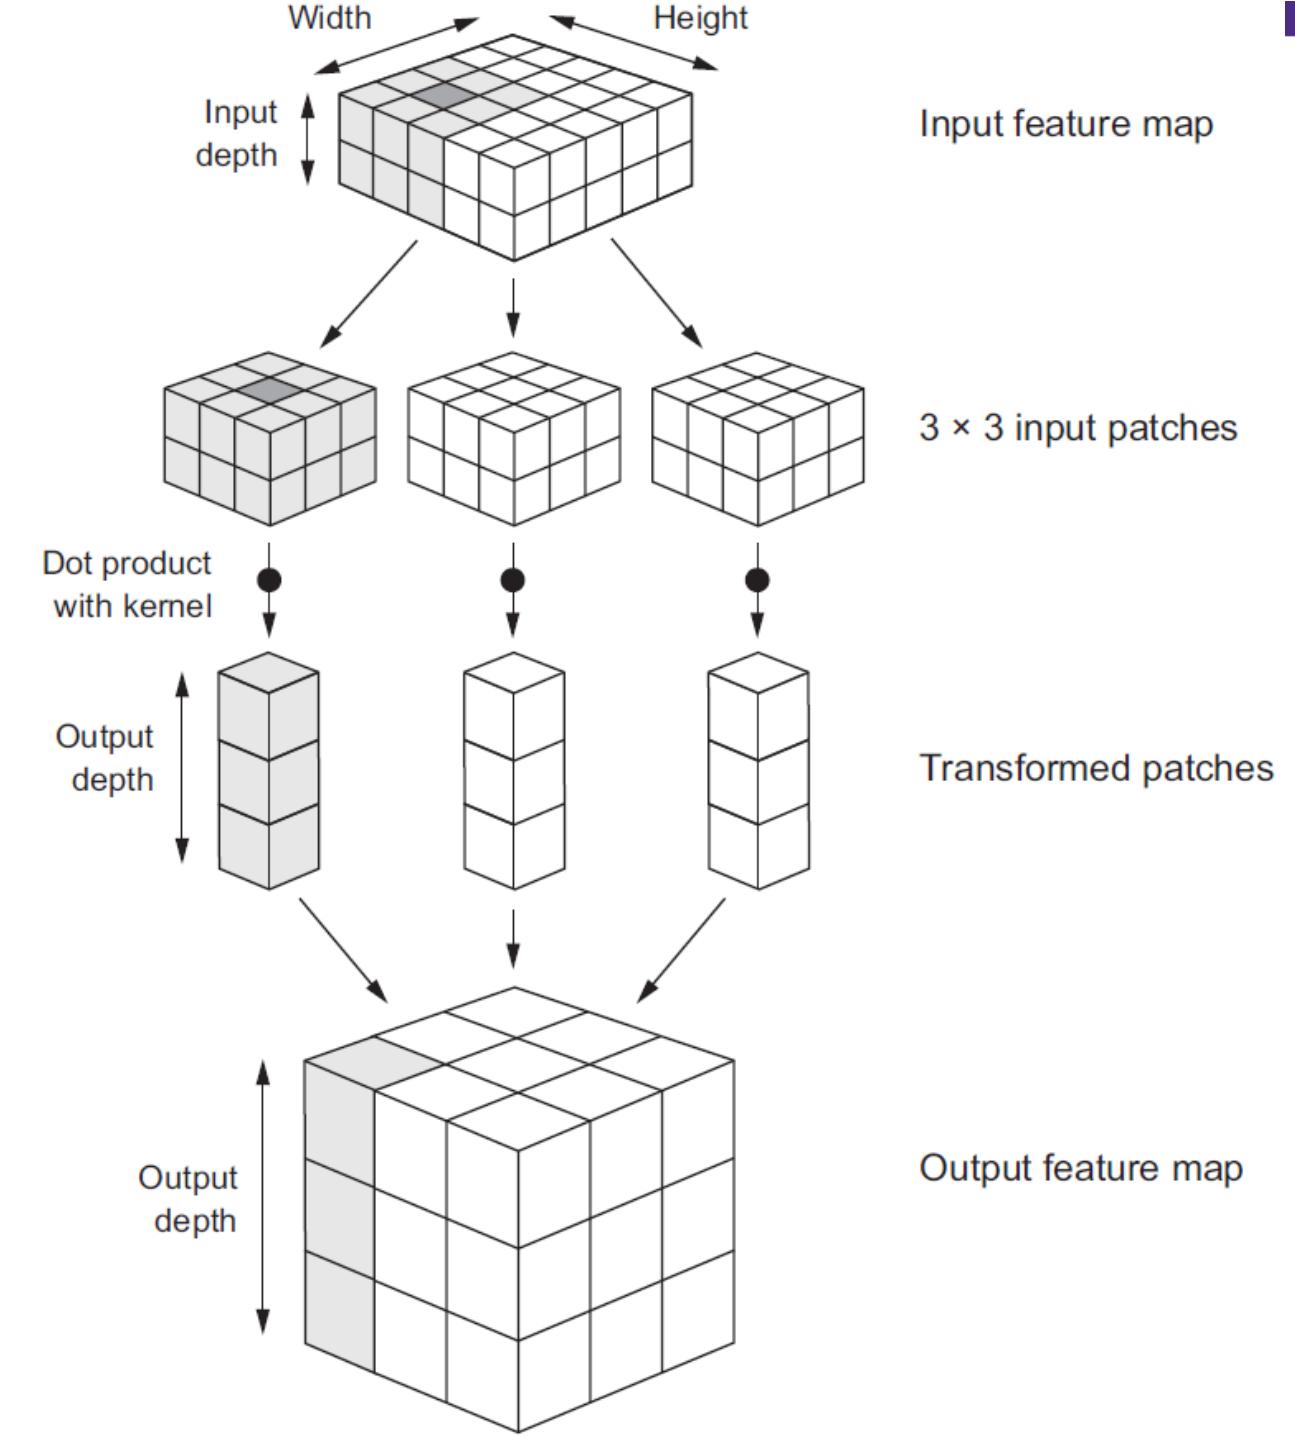


Key Parameters for Convolutions

- *Size of the patches extracted from the inputs*—These are typically 3×3 or 5×5 . In the example, they were 3×3 , which is a common choice.
- *Depth of the output feature map*—The number of filters computed by the convolution. The example started with a depth of 32 and ended with a depth of 64.



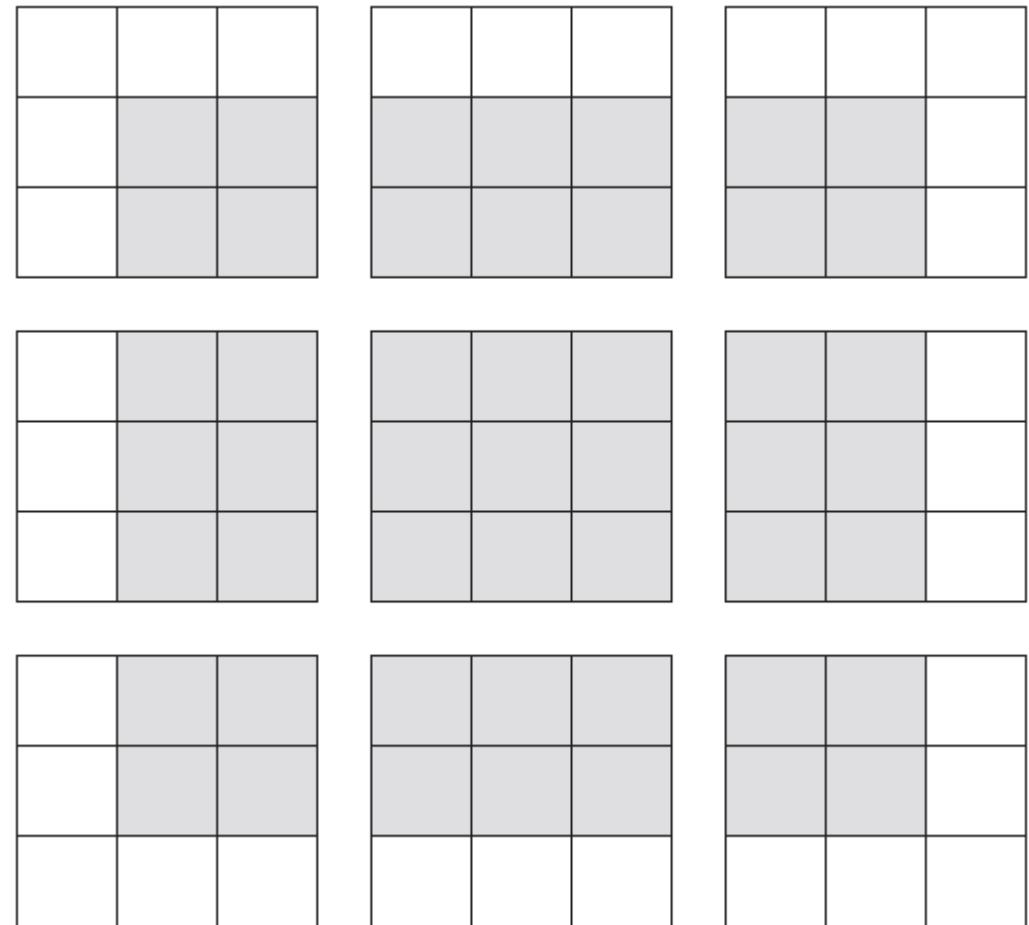
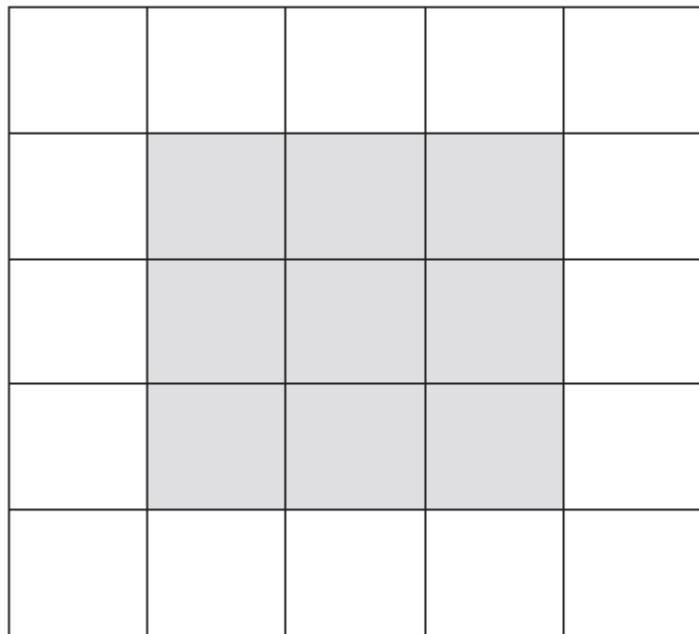
How Convolution Works



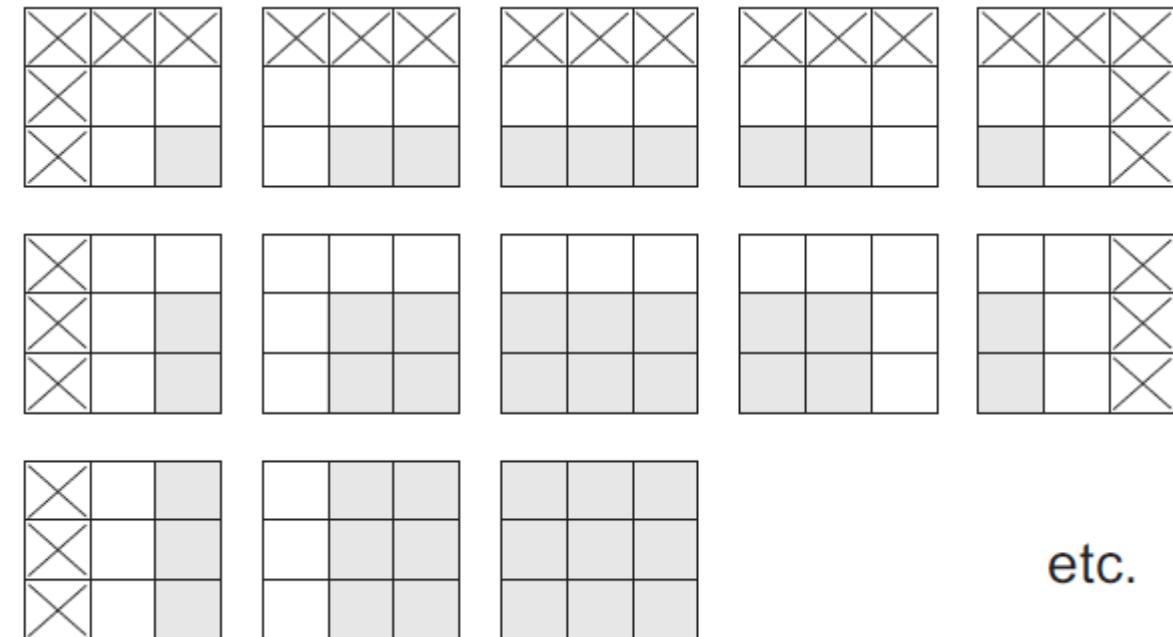
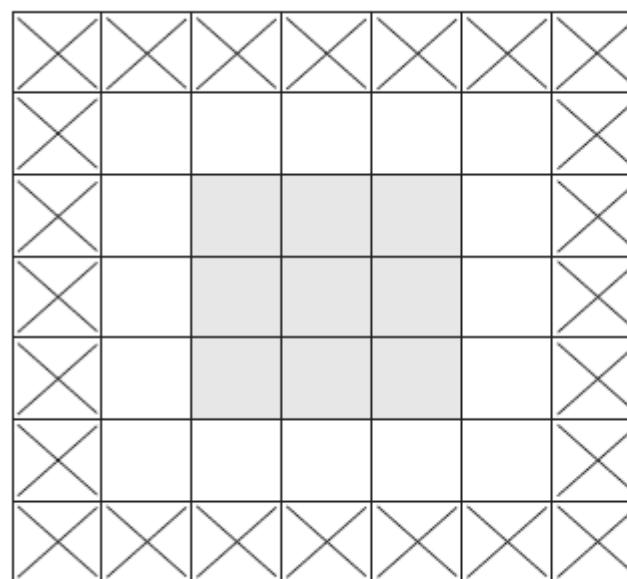
Output Width and Height May Differ from Input Width and Height

- Border effects, which can be countered by padding the input feature map
- The use of *strides*, which I'll define in a second

“Valid” Locations of 3x3 Convolution Patches for a 5x5 Image

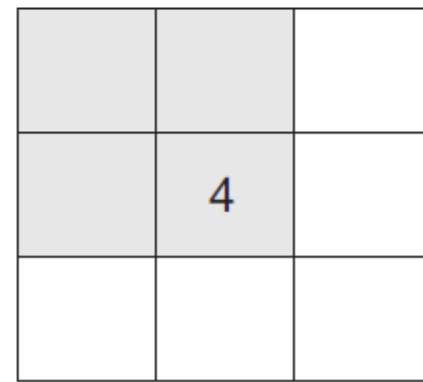
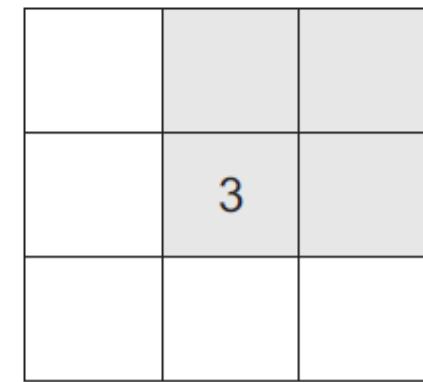
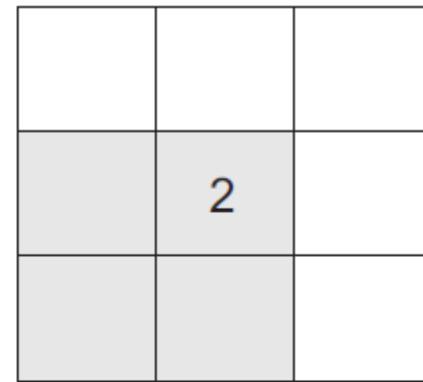
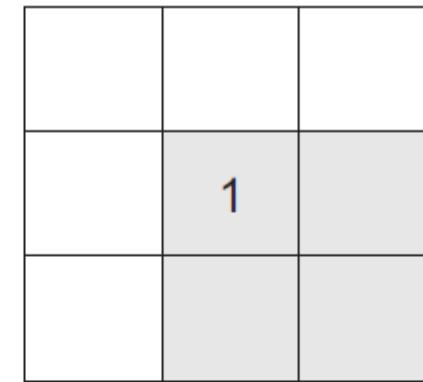
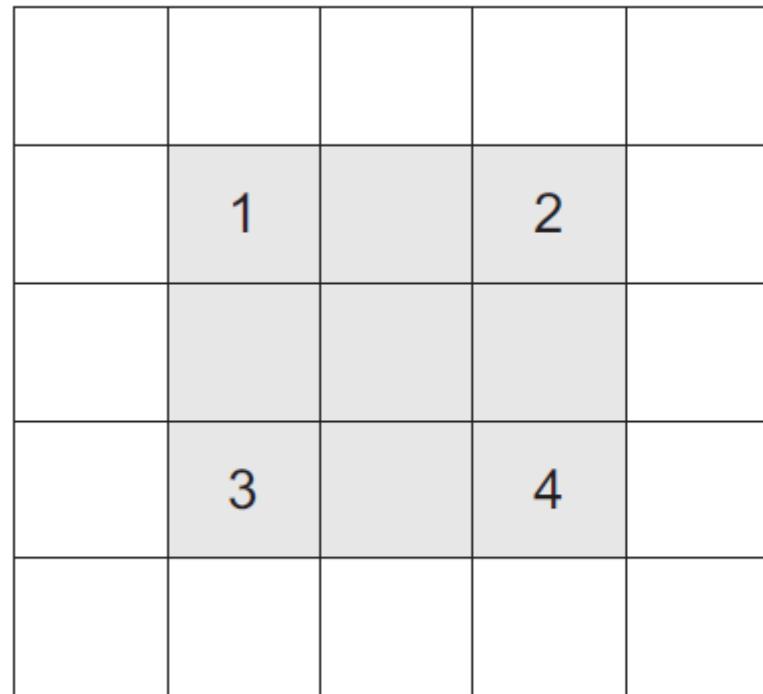


Padding a 5x5 Image for “Same” Size Output from a 3x3 Convolution Patch



etc.

3x3 Convolution Patches with 2x2 Strides



Model Without Max Pooling

```
model_no_max_pool = models.Sequential()
model_no_max_pool.add(layers.Conv2D(32, (3, 3), activation='relu',
                                   input_shape=(28, 28, 1)))
model_no_max_pool.add(layers.Conv2D(64, (3, 3), activation='relu'))
model_no_max_pool.add(layers.Conv2D(64, (3, 3), activation='relu'))

>>> model_no_max_pool.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 26, 26, 32)	320
<hr/>		
conv2d_5 (Conv2D)	(None, 24, 24, 64)	18496
<hr/>		
conv2d_6 (Conv2D)	(None, 22, 22, 64)	36928
=====		
Total params:	55,744	
Trainable params:	55,744	
Non-trainable params:	0	



What's Wrong with Model Without Max Pooling?

- It isn't conducive to learning a spatial hierarchy of features. The 3×3 windows in the third layer will only contain information coming from 7×7 windows in the initial input. The high-level patterns learned by the convnet will still be very small with regard to the initial input, which may not be enough to learn to classify digits (try recognizing a digit by only looking at it through windows that are 7×7 pixels!). We need the features from the last convolution layer to contain information about the totality of the input.
- The final feature map has $22 \times 22 \times 64 = 30,976$ total coefficients per sample. This is huge. If you were to flatten it to stick a Dense layer of size 512 on top, that layer would have 15.8 million parameters. This is far too large for such a small model and would result in intense overfitting.

Dogs versus Cats

Images are heterogeneous in size and appearance



Preprocessing: Creating Partition Directories

- 25,000 images per class: 543 megabytes compressed
- For this exercise: 1,000 per class trn; 500 per class val; and 500 per class tst

```
Path to the directory where the  
original dataset was uncompressed  
import os, shutil  
  
original_dataset_dir = '/Users/fchollet/Downloads/kaggle_original_data'  
  
base_dir = '/Users/fchollet/Downloads/cats_and_dogs_small'  
os.mkdir(base_dir)  
  
train_dir = os.path.join(base_dir, 'train')  
os.mkdir(train_dir)  
validation_dir = os.path.join(base_dir, 'validation')  
os.mkdir(validation_dir)  
test_dir = os.path.join(base_dir, 'test')  
os.mkdir(test_dir)
```

Directory where you'll store
your smaller dataset

Directories for
the training,
validation, and
test splits



Preprocessing: Creating Class Directories

```
train_cats_dir = os.path.join(train_dir, 'cats')  
os.mkdir(train_cats_dir) | Directory with  
                           training cat pictures  
  
train_dogs_dir = os.path.join(train_dir, 'dogs')  
os.mkdir(train_dogs_dir) | Directory with  
                           training dog pictures  
  
validation_cats_dir = os.path.join(validation_dir, 'cats')  
os.mkdir(validation_cats_dir) | Directory with  
                               validation cat pictures  
  
validation_dogs_dir = os.path.join(validation_dir, 'dogs')  
os.mkdir(validation_dogs_dir) | Directory with  
                               validation dog pictures  
  
test_cats_dir = os.path.join(test_dir, 'cats')  
os.mkdir(test_cats_dir) | Directory with test cat pictures  
  
test_dogs_dir = os.path.join(test_dir, 'dogs')  
os.mkdir(test_dogs_dir) | Directory with test dog pictures
```



Preprocessing: Copying Cat Files

```
fnames = ['cat.{}.jpg'.format(i) for i in range(1000)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(train_cats_dir, fname)  
    shutil.copyfile(src, dst)
```

Copies the first
1,000 cat images
to `train_cats_dir`

```
fnames = ['cat.{}.jpg'.format(i) for i in range(1000, 1500)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(validation_cats_dir, fname)  
    shutil.copyfile(src, dst)
```

Copies the next 500
cat images to
`validation_cats_dir`

```
fnames = ['cat.{}.jpg'.format(i) for i in range(1500, 2000)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(test_cats_dir, fname)  
    shutil.copyfile(src, dst)
```

Copies the next 500
cat images to
`test_cats_dir`



Preprocessing: Copying Dog Files

```
fnames = ['dog.{}.jpg'.format(i) for i in range(1000)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(train_dogs_dir, fname)  
    shutil.copyfile(src, dst)
```

Copies the first
1,000 dog images
to train_dogs_dir

```
fnames = ['dog.{}.jpg'.format(i) for i in range(1000, 1500)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(validation_dogs_dir, fname)  
    shutil.copyfile(src, dst)
```

Copies the next 500
dog images to
validation_dogs_dir

```
fnames = ['dog.{}.jpg'.format(i) for i in range(1500, 2000)]  
for fname in fnames:  
    src = os.path.join(original_dataset_dir, fname)  
    dst = os.path.join(test_dogs_dir, fname)  
    shutil.copyfile(src, dst)
```

Copies the next 500
dog images to
test_dogs_dir



Sanity Check: File Counts

```
>>> print('total training cat images:', len(os.listdir(train_cats_dir)))
total training cat images: 1000
>>> print('total training dog images:', len(os.listdir(train_dogs_dir)))
total training dog images: 1000
>>> print('total validation cat images:', len(os.listdir(validation_cats_dir)))
total validation cat images: 500
>>> print('total validation dog images:', len(os.listdir(validation_dogs_dir)))
total validation dog images: 500
>>> print('total test cat images:', len(os.listdir(test_cats_dir)))
total test cat images: 500
>>> print('total test dog images:', len(os.listdir(test_dogs_dir)))
total test dog images: 500
```



ConvNet for Dogs versus Cats: Definition

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                      input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

ConvNet for Dogs versus Cats: Summary

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
maxpooling2d_1 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_2 (Conv2D)	(None, 72, 72, 64)	18496
maxpooling2d_2 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
maxpooling2d_3 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
maxpooling2d_4 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 512)	3211776
dense_2 (Dense)	(None, 1)	513
<hr/>		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		



Configuring the Model

```
from keras import optimizers  
  
model.compile(loss='binary_crossentropy',  
                optimizer=optimizers.RMSprop(lr=1e-4),  
                metrics=['acc'])
```

Steps for Creating Tensors

- 1 Read the picture files.
- 2 Decode the JPEG content to RGB grids of pixels.
- 3 Convert these into floating-point tensors.
- 4 Rescale the pixel values (between 0 and 255) to the [0, 1] interval (as you know, neural networks prefer to deal with small input values).

Code for Creating Tensors

```
from keras.preprocessing.image import ImageDataGenerator  
  
train_datagen = ImageDataGenerator(rescale=1./255) | Rescales all images by 1/255  
test_datagen = ImageDataGenerator(rescale=1./255)  
  
train_generator = train_datagen.flow_from_directory(  
    train_dir,  
    target_size=(150, 150)  ← Resizes all images to 150 × 150  
    batch_size=20,  
    class_mode='binary')  
  
validation_generator = test_datagen.flow_from_directory(  
    validation_dir,  
    target_size=(150, 150),  
    batch_size=20,  
    class_mode='binary')  
  
Target directory
```

Because you use
`binary_crossentropy`
loss, you need binary
labels.

Understanding Python Generators

```
def generator():
    i = 0
    while True:
        i += 1
        yield i

for item in generator():
    print(item)
    if item > 4:
        break

1
2
3
4
5
```



Example Train Batch

```
>>> for data_batch, labels_batch in train_generator:  
>>>     print('data batch shape:', data_batch.shape)  
>>>     print('labels batch shape:', labels_batch.shape)  
>>>     break  
data batch shape: (20, 150, 150, 3)  
labels batch shape: (20, )
```



Fitting and Saving the Model

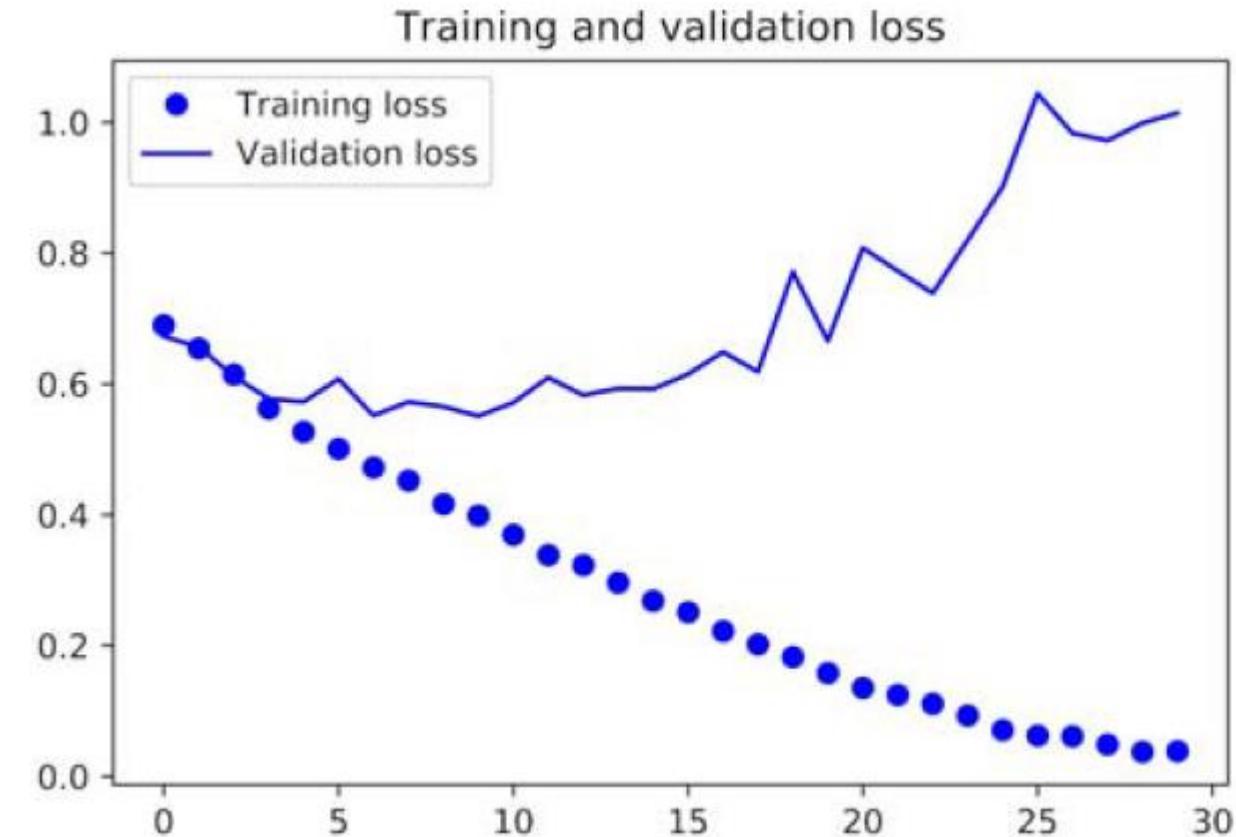
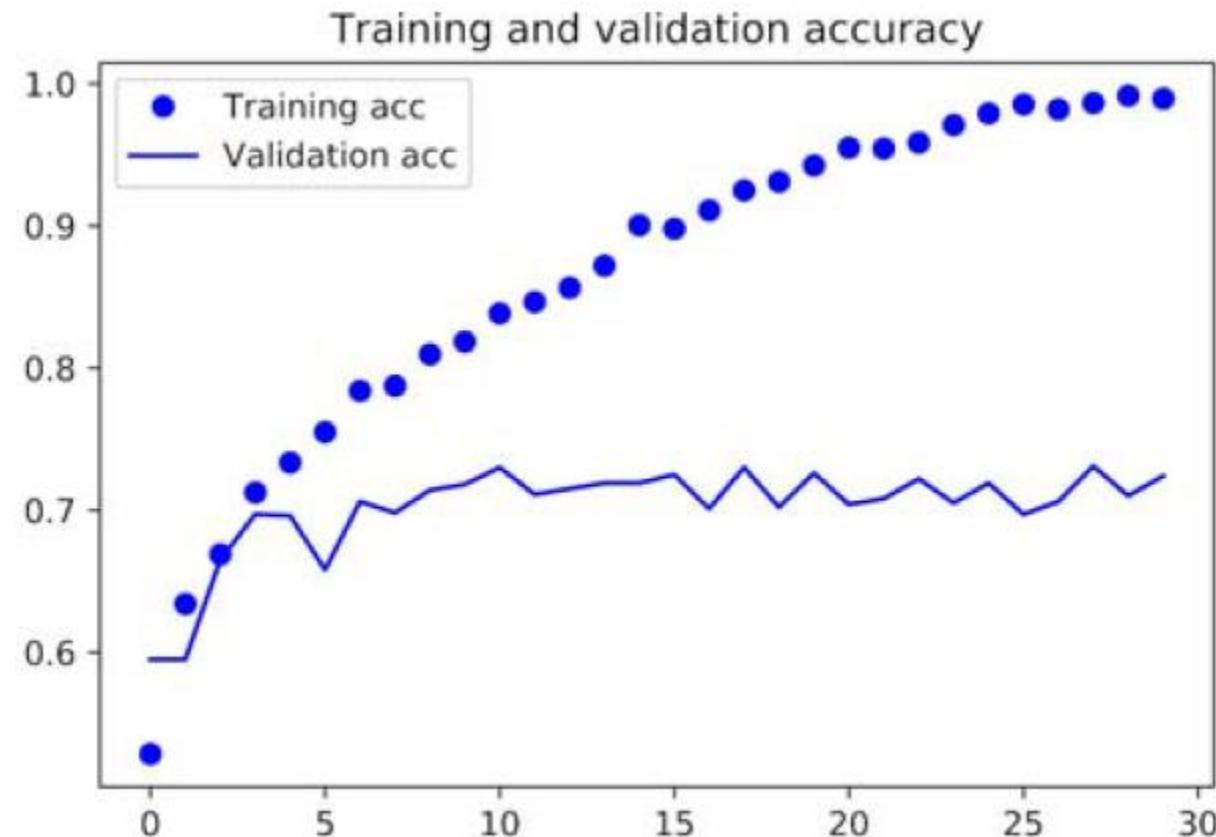
```
steps_per_epoch = trnX.shape[0] / batch_size
```

```
validation_steps = valX.shape[0] / batch_size
```

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=30,  
    validation_data=validation_generator,  
    validation_steps=50)
```

```
model.save('cats_and_dogs_small_1.h5')
```

Accuracy and Cross Entropy

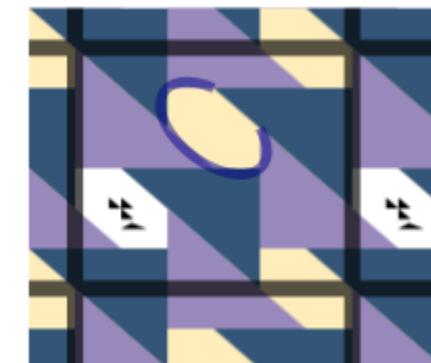


Training new model with small data: 70-72% accuracy

Data Augmentation

- Fill modes
 - “constant”: kkkkkkkk|abcd|kkkkkkkk (cval=k)
 - “nearest”: aaaaaaaaa|abcd|ddddddddd [extend “last” pixel to border]
 - “reflect”: abcddcba|abcd|dcbaabcd
 - “wrap”: abcdabcd|abcd|abcdabcd

```
datagen = ImageDataGenerator(  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')
```





Data Augmentation Options

- `rotation_range` is a value in degrees (0–180), a range within which to randomly rotate pictures.
- `width_shift` and `height_shift` are ranges (as a fraction of total width or height) within which to randomly translate pictures vertically or horizontally.
- `shear_range` is for randomly applying shearing transformations.
- `zoom_range` is for randomly zooming inside pictures.
- `horizontal_flip` is for randomly flipping half the images horizontally—relevant when there are no assumptions of horizontal asymmetry (for example, real-world pictures).
- `fill_mode` is the strategy used for filling in newly created pixels, which can appear after a rotation or a width/height shift.

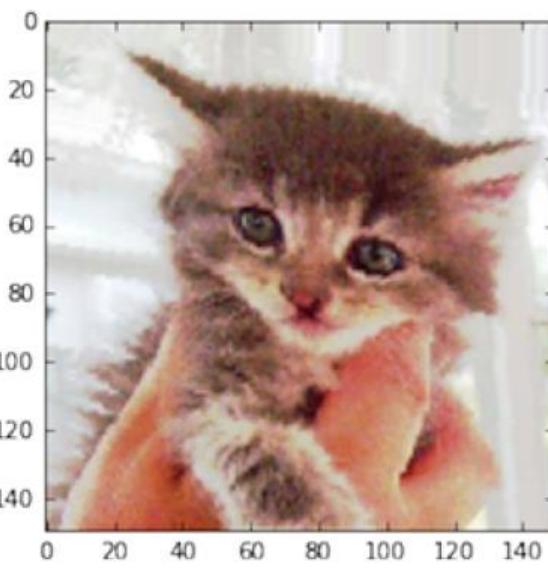
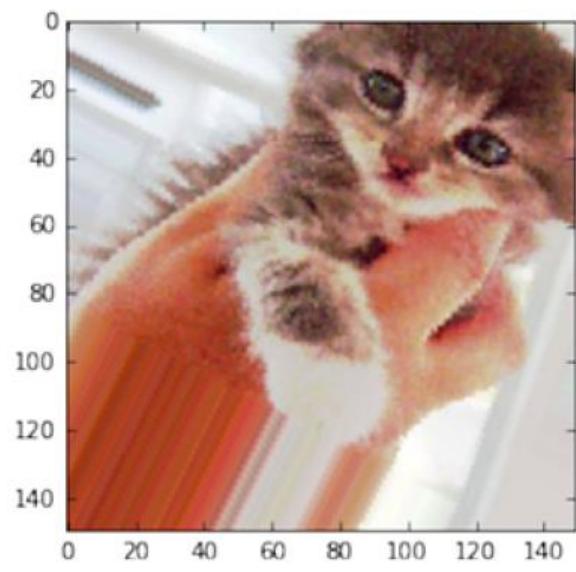
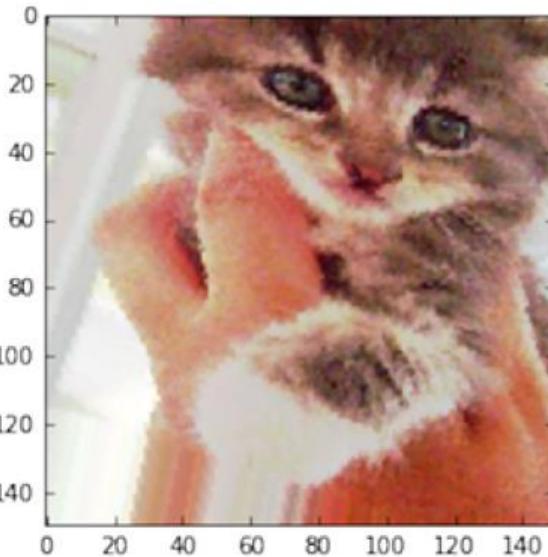
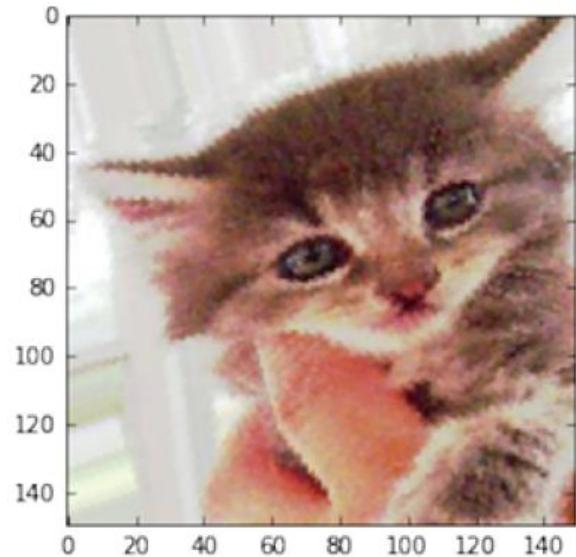
Displaying Augmented Training Images

```
from keras.preprocessing import image           ← Module with image-  
fnames = [os.path.join(train_cats_dir, fname) for preprocessing utilities  
    fname in os.listdir(train_cats_dir)]  
  
img_path = fnames[3]                          ← Chooses one image to augment  
  
img = image.load_img(img_path, target_size=(150, 150)) ← Reads the image  
                                                        and resizes it  
  
x = image.img_to_array(img)      ← Converts it to a Numpy array with shape (150, 150, 3)  
  
x = x.reshape((1,) + x.shape)        ← Reshapes it to (1, 150, 150, 3)  
  
i = 0  
for batch in datagen.flow(x, batch_size=1):  
    plt.figure(i)  
    imgplot = plt.imshow(image.array_to_img(batch[0]))  
    i += 1  
    if i % 4 == 0:  
        break  
  
plt.show()
```

Generates batches of randomly transformed images. Loops indefinitely, so you need to break the loop at some point!

Augmented Images

- cat.100.jpg: 499x403 => 150x150
- Transformations include Horizontal Flip, Rotation, Shear, Translation, and Zoom





Model with Dropout for Augmented Data

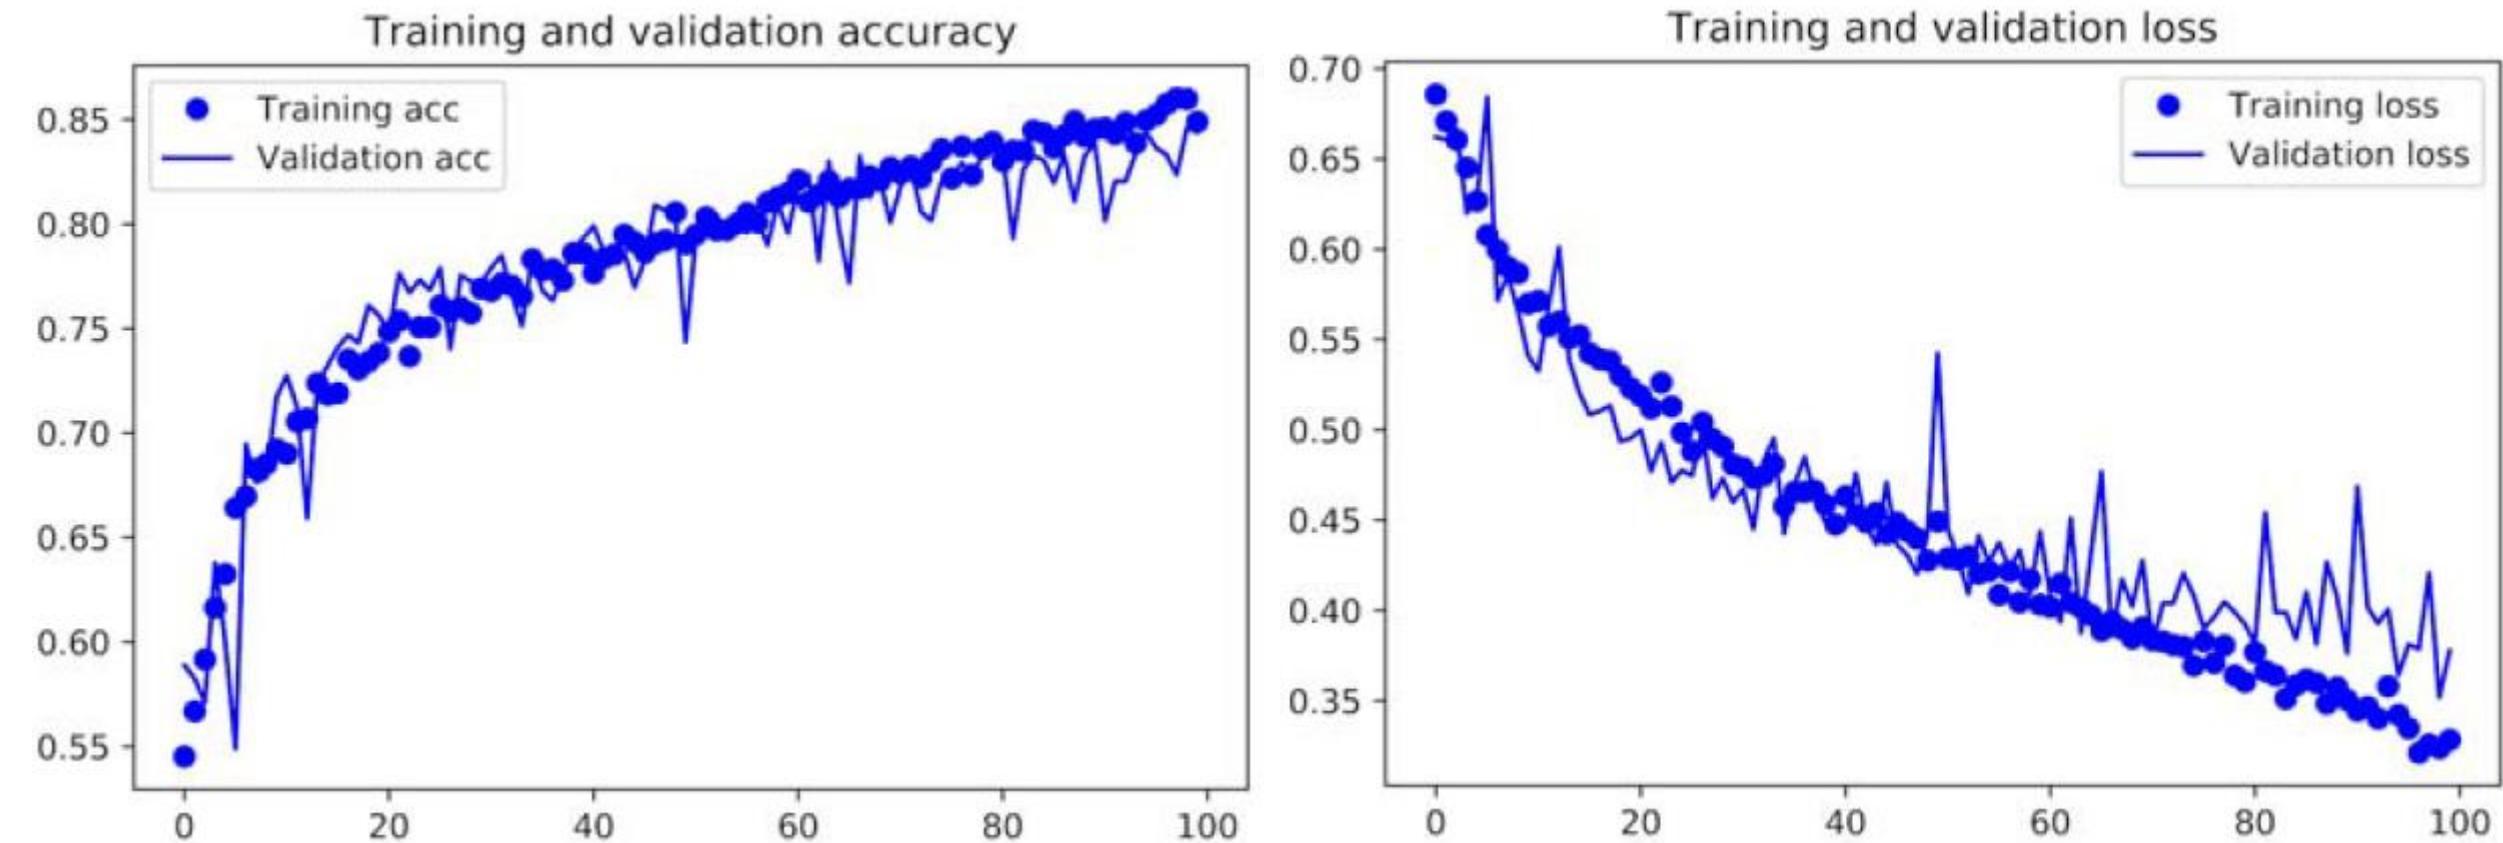
```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                      input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=['acc'])
```

Training the ConvNet with Augmented Data

```
train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,)  
  
test_datagen = ImageDataGenerator(rescale=1./255) ← Note that the validation data shouldn't be augmented!  
  
train_generator = train_datagen.flow_from_directory(  
    train_dir,  
    target_size=(150, 150), ← Resizes all images to 150 × 150  
    batch_size=32,  
    class_mode='binary') ← Because you use binary_crossentropy loss, you need binary labels.  
  
validation_generator = test_datagen.flow_from_directory(  
    validation_dir,  
    target_size=(150, 150),  
    batch_size=32,  
    class_mode='binary')  
  
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=100,  
    epochs=100,  
    validation_data=validation_generator,  
    validation_steps=50)
```

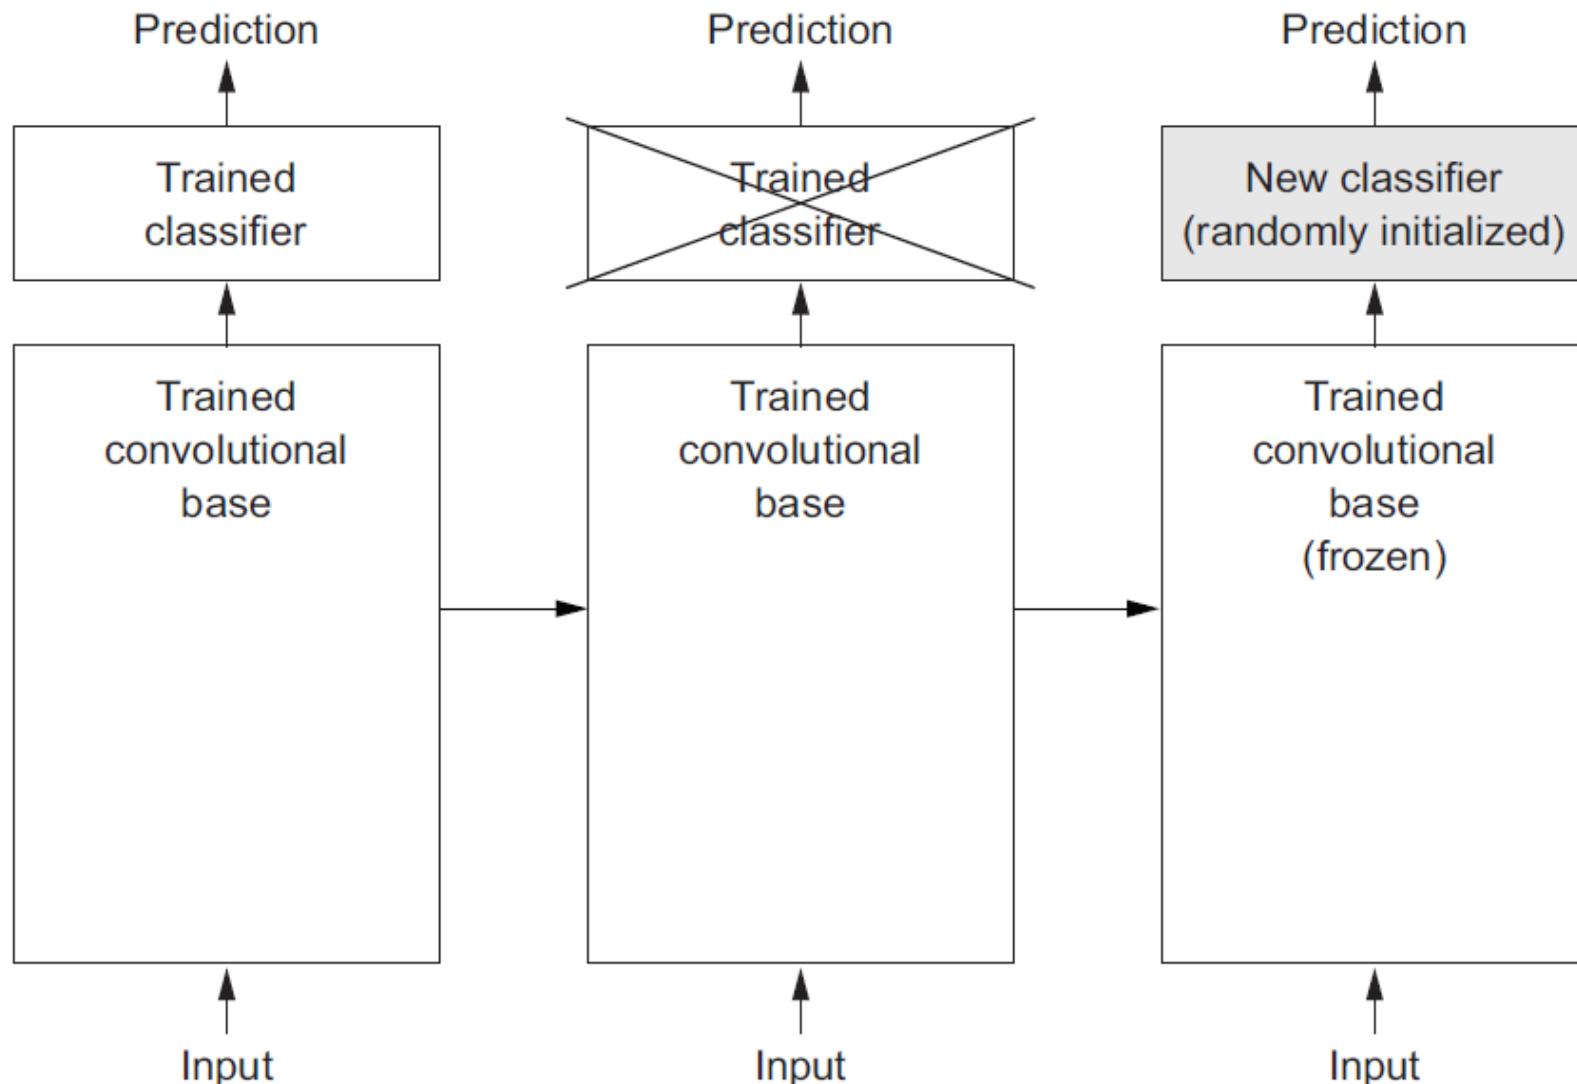
Accuracy and Cross Entropy



Training new model with small data (and augmentation and dropout): 82% accuracy [71% to 82%: +15% relative improvement]

Swapping Classifiers for the Same Convolutional Base

- Transferring learning from one application to another
- Visual features from one application can help in another application, especially when there is less data for the new application





VGG16

- ConvNet built by the University of Oxford's Visual Geometry Group
- 16 Layers have Parameters (weights and biases)

```
from keras.applications import VGG16

conv_base = VGG16(weights='imagenet',
                  include_top=False,
                  input_shape=(150, 150, 3))
```

- weights specifies the weight checkpoint from which to initialize the model.
- include_top refers to including (or not) the densely connected classifier on top of the network. By default, this densely connected classifier corresponds to the 1,000 classes from ImageNet. Because you intend to use your own densely connected classifier (with only two classes: cat and dog), you don't need to include it.
- input_shape is the shape of the image tensors that you'll feed to the network. This argument is purely optional: if you don't pass it, the network will be able to process inputs of any size.

VGG16 Summary: Blocks 1 and 2

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Convolution2D)	(None, 150, 150, 64)	1792
block1_conv2 (Convolution2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Convolution2D)	(None, 75, 75, 128)	73856
block2_conv2 (Convolution2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0



VGG16 Summary: Blocks 3 and 4

block3_conv1 (Convolution2D)	(None, 37, 37, 256)	295168
block3_conv2 (Convolution2D)	(None, 37, 37, 256)	590080
block3_conv3 (Convolution2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Convolution2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Convolution2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Convolution2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0



VGG16 Summary: Block 5

block5_conv1 (Convolution2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Convolution2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Convolution2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params:	14,714,688	
Trainable params:	14,714,688	
Non-trainable params:	0	



VGG16 Summary: Top [removed for this application]

3 fully connected layers ...

flatten (Flatten)	(None, 25088)	0
<hr/>		
fc1 (Dense)	(None, 4096)	102,764,544
<hr/>		
fc2 (Dense)	(None, 4096)	16,781,312
<hr/>		
predictions (Dense)	(None, 1000)	4,097,000



Alternatives for Using the Convolutional Base

- Running the convolutional base over your dataset, recording its output to a Numpy array on disk, and then using this data as input to a standalone, densely connected classifier similar to those you saw in part 1 of this book. This solution is fast and cheap to run, because it only requires running the convolutional base once for every input image, and the convolutional base is by far the most expensive part of the pipeline. But for the same reason, this technique won't allow you to use data augmentation.
- Extending the model you have (`conv_base`) by adding Dense layers on top, and running the whole thing end to end on the input data. This will allow you to use data augmentation, because every input image goes through the convolutional base every time it's seen by the model. But for the same reason, this technique is far more expensive than the first.

Extracting Features: Define the Image Generator

```
import os
import numpy as np
from keras.preprocessing.image import ImageDataGenerator

base_dir = '/Users/fchollet/Downloads/cats_and_dogs_small'
train_dir = os.path.join(base_dir, 'train')
validation_dir = os.path.join(base_dir, 'validation')
test_dir = os.path.join(base_dir, 'test')

datagen = ImageDataGenerator(rescale=1./255)
batch_size = 20
```

Extracting Features: for Trn, Val, and Tst

```
def extract_features(directory, sample_count):
    features = np.zeros(shape=(sample_count, 4, 4, 512))
    labels = np.zeros(shape=(sample_count))
    generator = datagen.flow_from_directory(
        directory,
        target_size=(150, 150),
        batch_size=batch_size,
        class_mode='binary')
    i = 0
    for inputs_batch, labels_batch in generator:
        features[i * batch_size : (i + 1) * batch_size] = features_batch
        labels[i * batch_size : (i + 1) * batch_size] = labels_batch
        i += 1
        if i * batch_size >= sample_count:
            break
    return features, labels

train_features, train_labels = extract_features(train_dir, 2000)
validation_features, validation_labels = extract_features(validation_dir, 1000)
test_features, test_labels = extract_features(test_dir, 1000)
```

Note that because generators yield data indefinitely in a loop, you must break after every image has been seen once.



Training a Model with the Convolutional Features

```
train_features = np.reshape(train_features, (2000, 4 * 4 * 512))
validation_features = np.reshape(validation_features, (1000, 4 * 4 * 512))
test_features = np.reshape(test_features, (1000, 4 * 4 * 512))

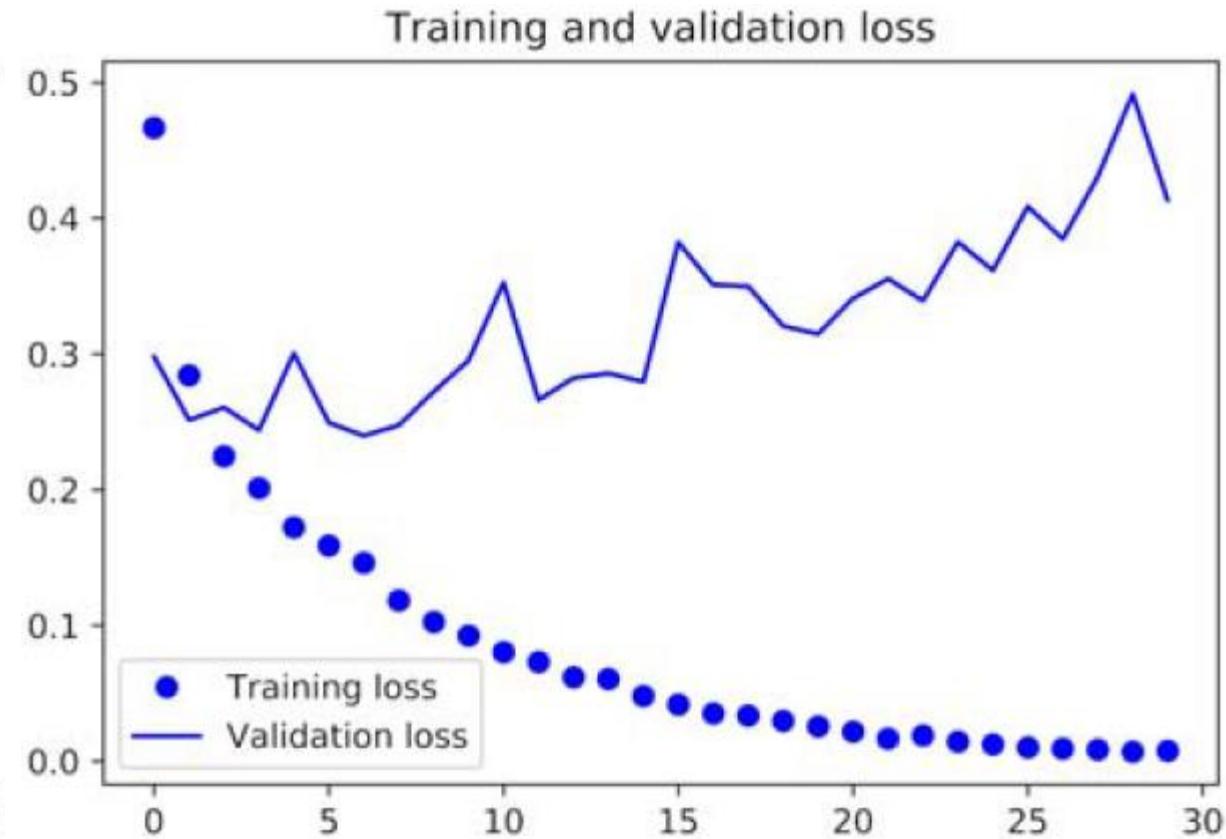
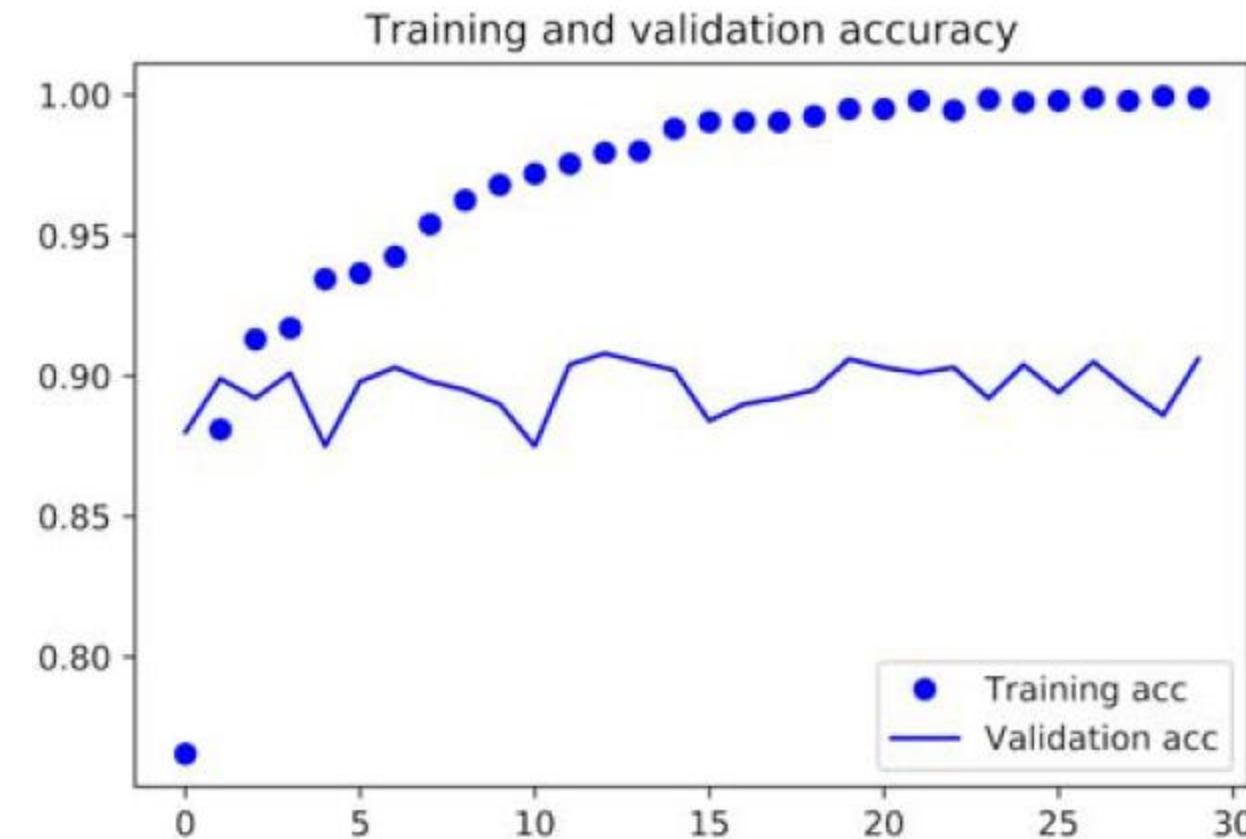
from keras import models
from keras import layers
from keras import optimizers

model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_dim=4 * 4 * 512))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer=optimizers.RMSprop(lr=2e-5),
              loss='binary_crossentropy',
              metrics=['acc'])

history = model.fit(train_features, train_labels,
                     epochs=30,
                     batch_size=20,
                     validation_data=(validation_features, validation_labels))
```

Accuracy and Cross Entropy



Training new model with VGG16 features: 90% accuracy



Feature Extraction with Data Augmentation: the Warning

NOTE This technique is so expensive that you should only attempt it if you have access to a GPU—it's absolutely intractable on CPU. If you can't run your code on GPU, then the previous technique is the way to go.

Using the Convolutional Base as a Layer

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

>>> model.summary()

Layer (type)                  Output Shape        Param #
=====
vgg16 (Model)                (None, 4, 4, 512)    14714688
=====
flatten_1 (Flatten)           (None, 8192)         0
=====
dense_1 (Dense)               (None, 256)          2097408
=====
dense_2 (Dense)               (None, 1)             257
=====

Total params: 16,812,353
Trainable params: 16,812,353
Non-trainable params: 0
```



Freezing the Convolutional Base

$$(13 + 2) * 2 = 30$$

$$(0 + 2) * 2 = 4$$

2: Weight Matrix plus Bias Vector

```
>>> print('This is the number of trainable weights '
       'before freezing the conv base:', len(model.trainable_weights))
This is the number of trainable weights before freezing the conv base: 30
>>> conv_base.trainable = False
>>> print('This is the number of trainable weights '
       'after freezing the conv base:', len(model.trainable_weights))
This is the number of trainable weights after freezing the conv base: 4
```

Configuring the Data Generators

```
from keras.preprocessing.image import ImageDataGenerator  
from keras import optimizers  
  
train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')  
  
test_datagen = ImageDataGenerator(rescale=1./255)  
  
train_generator = train_datagen.flow_from_directory(  
    target_directory,          ← Target directory  
    train_dir,  
    target_size=(150, 150),    ← Resizes all images to 150 × 150  
    batch_size=20,  
    class_mode='binary')  
  
validation_generator = test_datagen.flow_from_directory(  
    validation_dir,  
    target_size=(150, 150),  
    batch_size=20,  
    class_mode='binary')
```

Note that the validation data shouldn't be augmented!

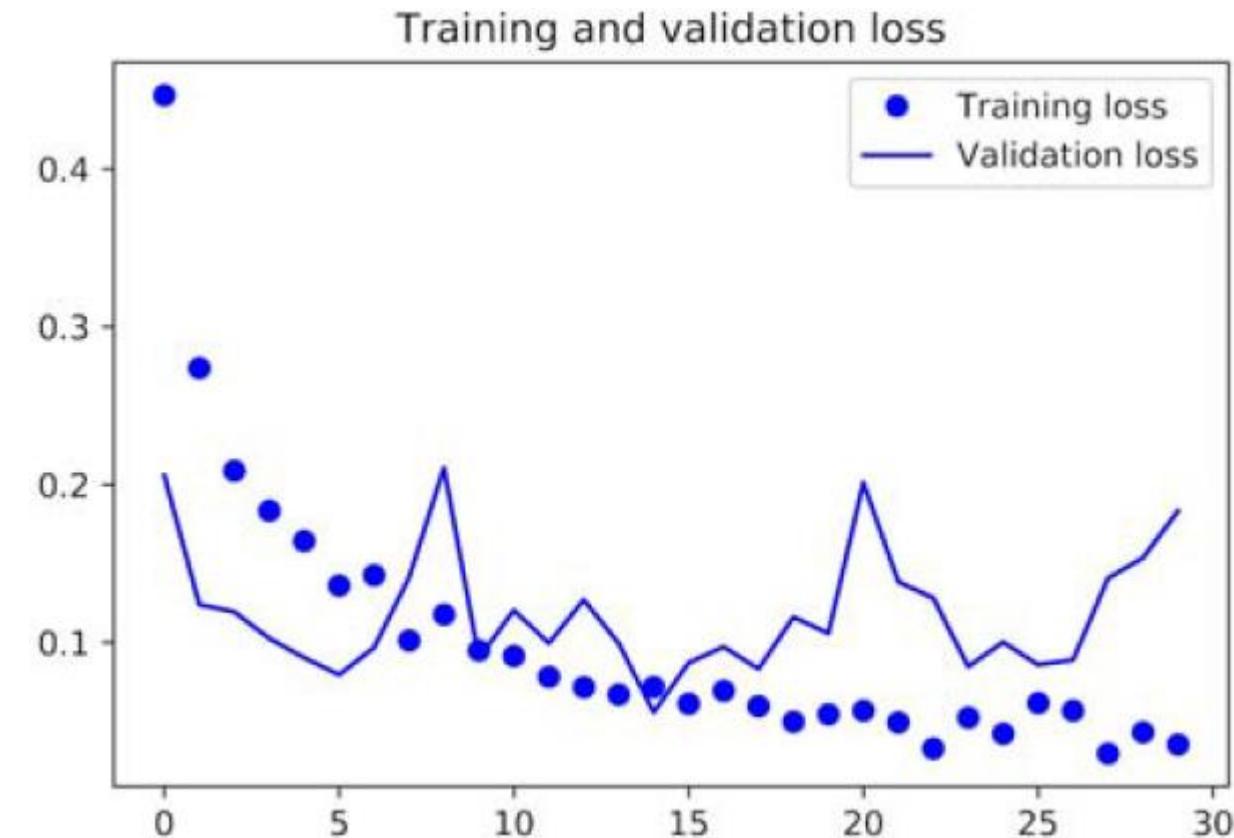
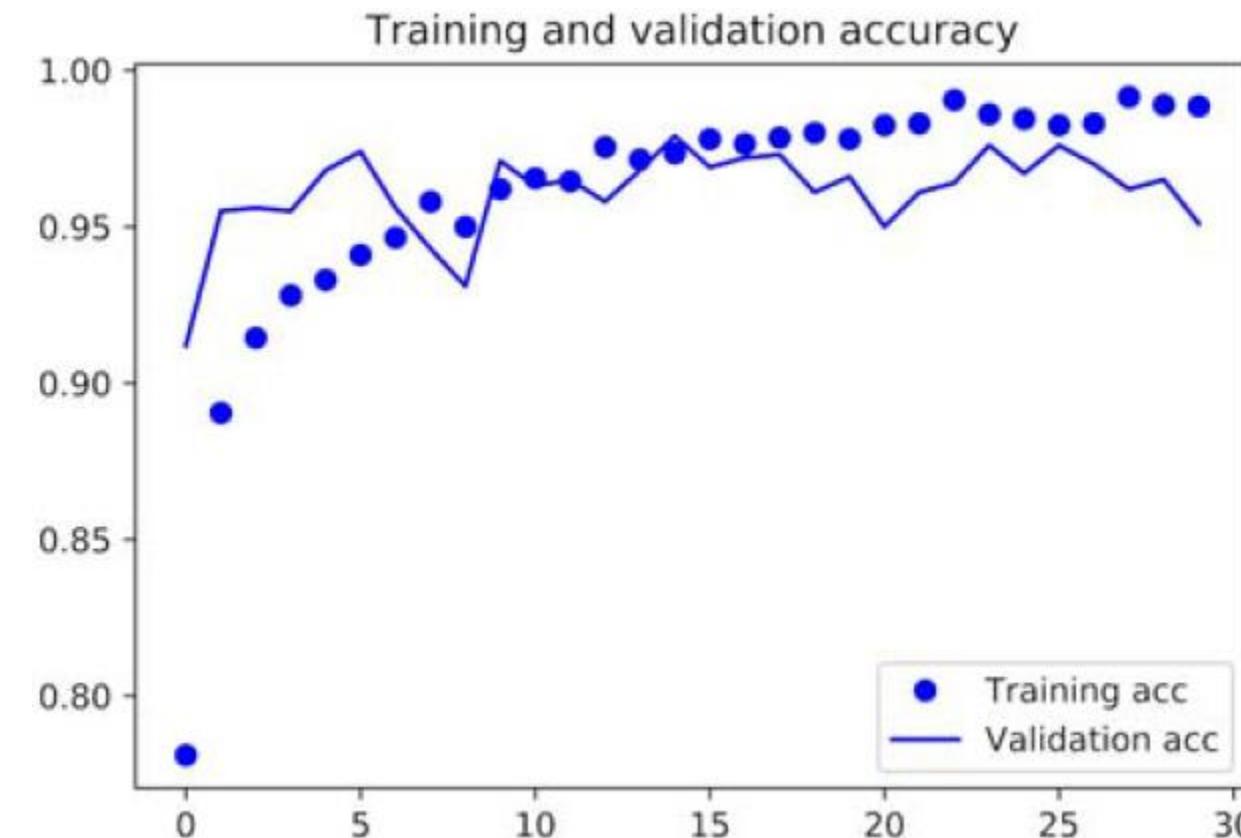
Because you use **binary_crossentropy** loss, you need binary labels.

Running the Generators with the Pretrained Base

```
model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=2e-5),
              metrics=['acc'])

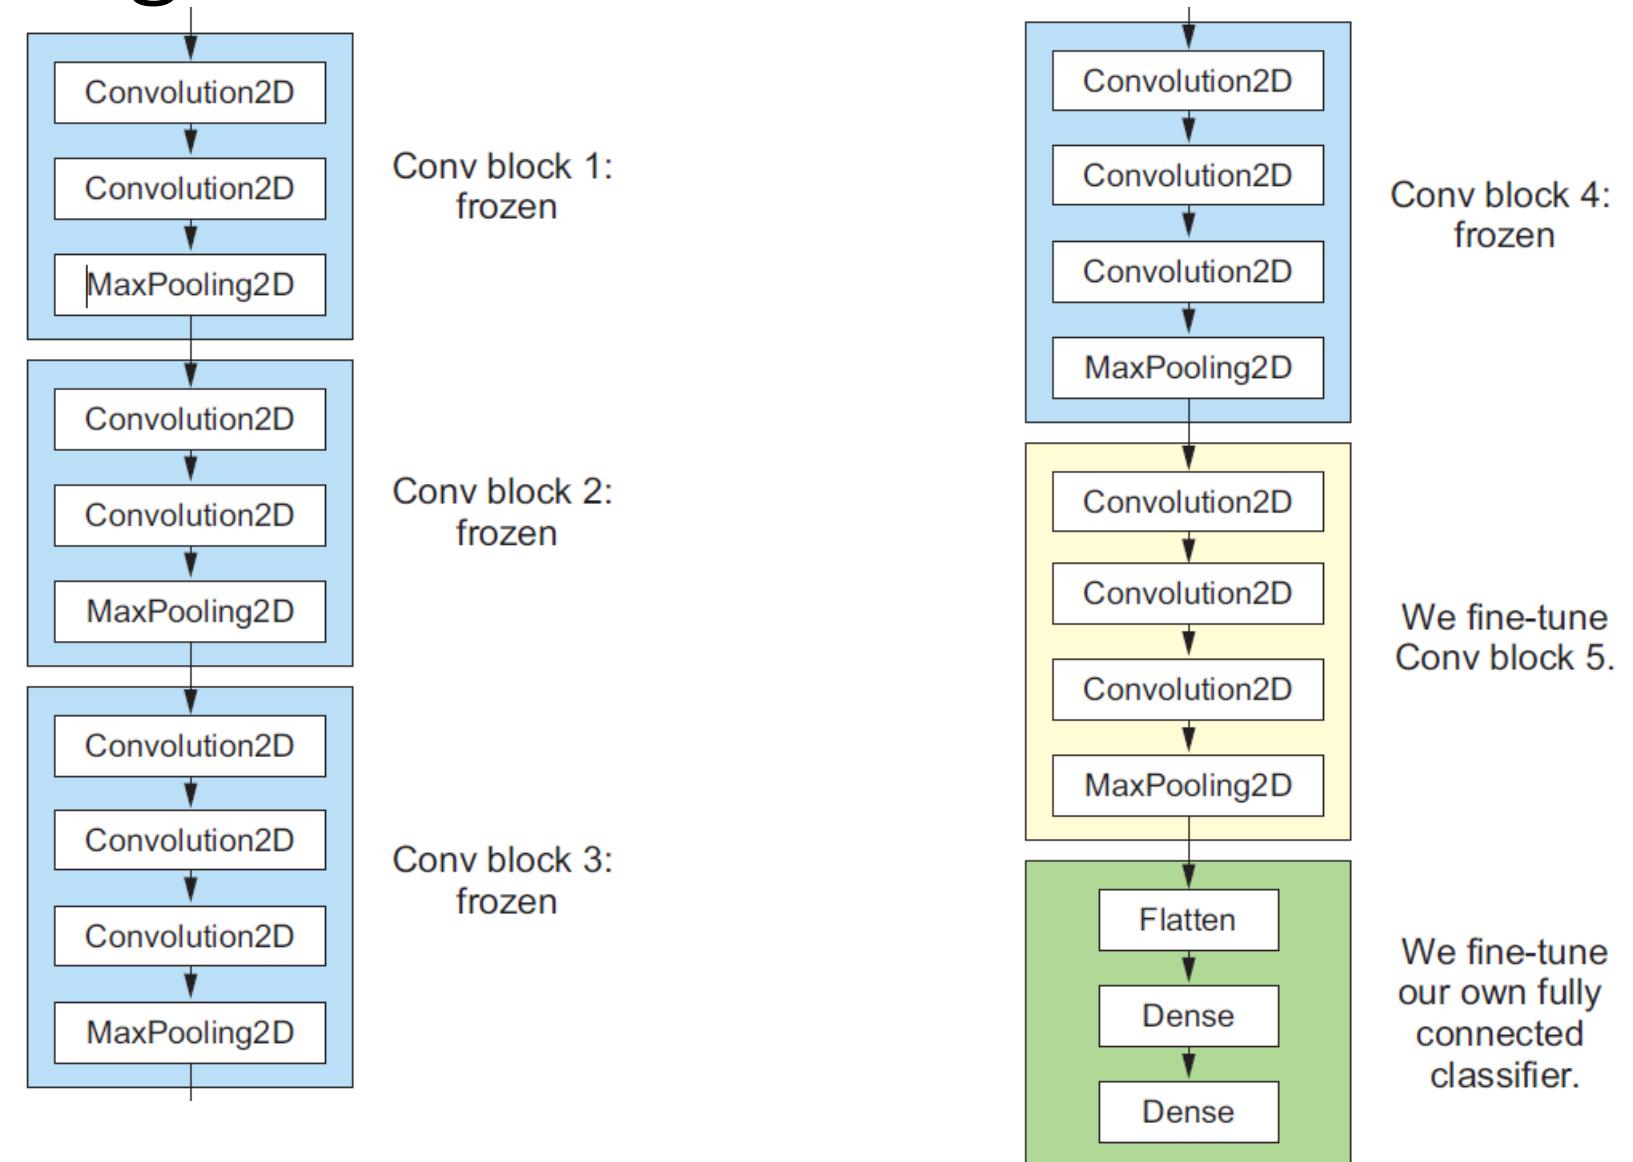
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=30,
    validation_data=validation_generator,
    validation_steps=50)
```

Accuracy and Cross Entropy



Training new model with VGG16 features, with data augmentation: 96% accuracy

Fine-Tuning Last Block of the VGG16 Base





Fine-Tuning a Pretrained Network

- 1 Add your custom network on top of an already-trained base network.
- 2 Freeze the base network.
- 3 Train the part you added.
- 4 Unfreeze some layers in the base network.
- 5 Jointly train both these layers and the part you added.

```
conv_base.trainable = True

set_trainable = False
for layer in conv_base.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable:
        layer.trainable = True
    else:
        layer.trainable = False
```



Fine-Tuning the Last Block of the VGG16 Base

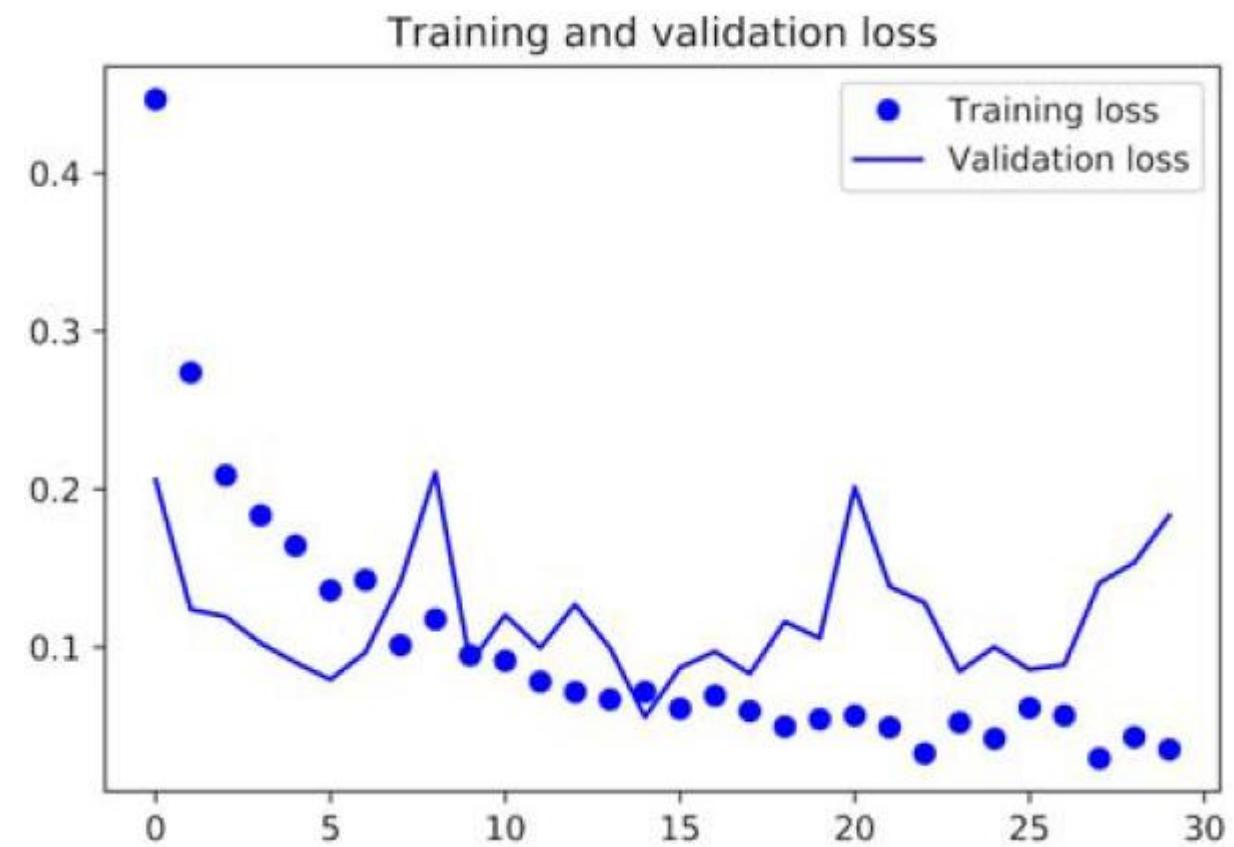
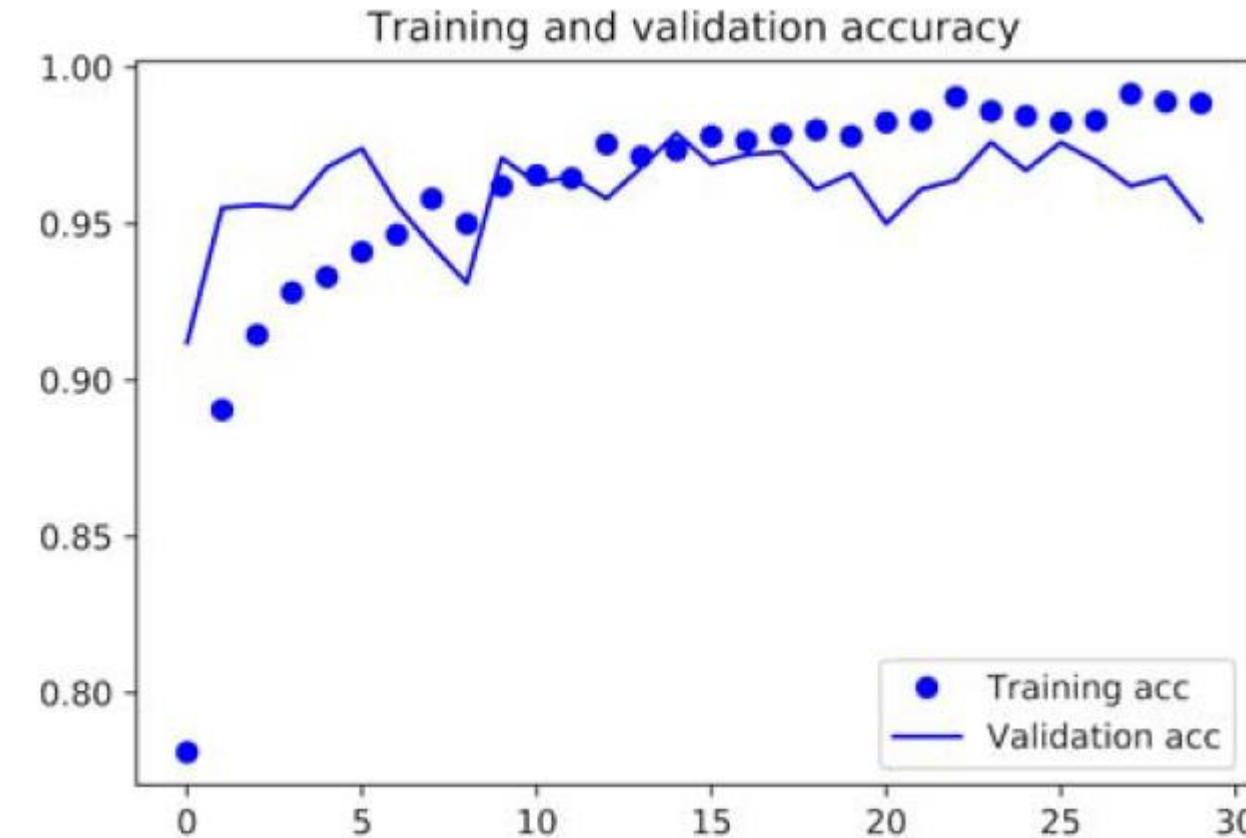
Half the learning rate

More epochs

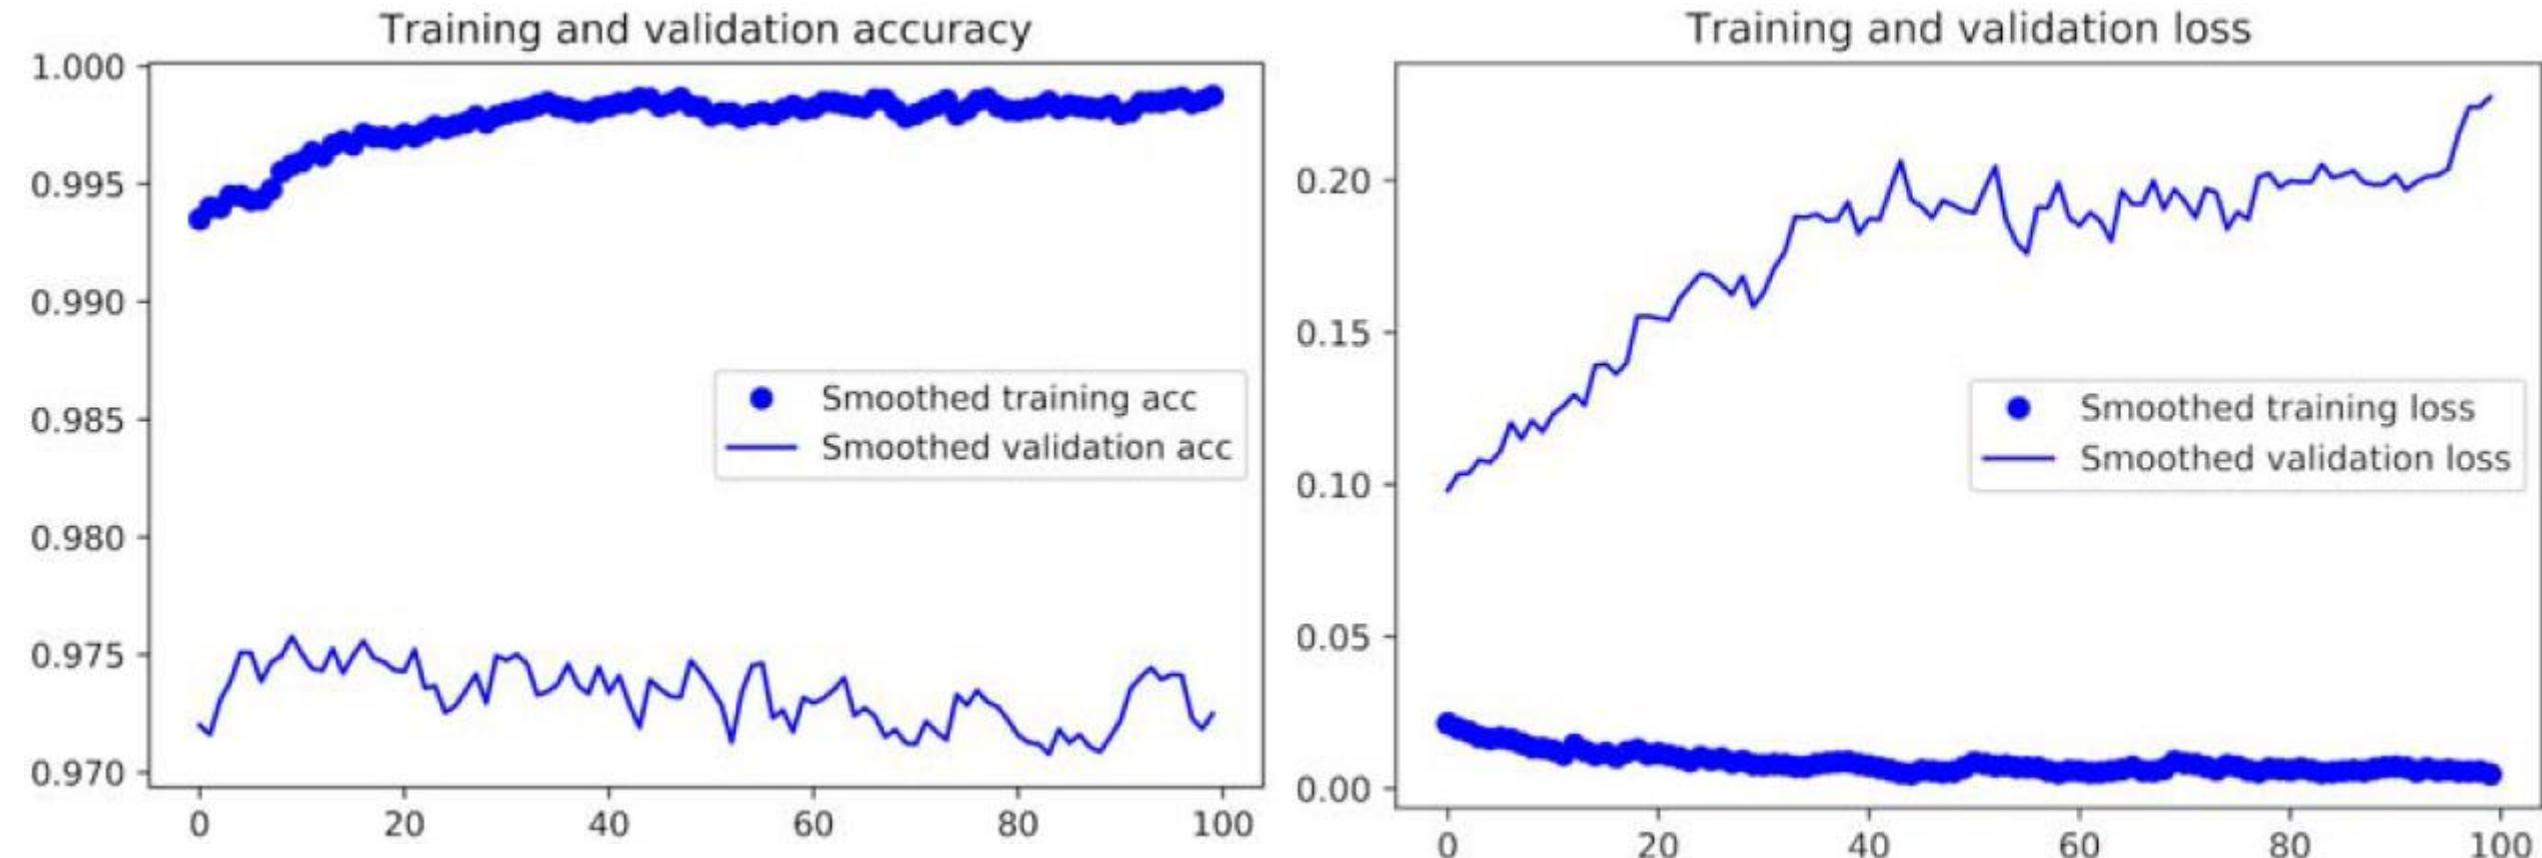
```
model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-5),
              metrics=['acc'])

history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=100,
    validation_data=validation_generator,
    validation_steps=50)
```

Accuracy and Cross Entropy



Smoothed Accuracy and Cross Entropy



Fine-tuning last block of VGG16 base, with data augmentation: 97% accuracy ... finally ready for testing!



Wrapping Up

- Convnets are the best type of machine-learning models for computer-vision tasks. It's possible to train one from scratch even on a very small dataset, with decent results.
- On a small dataset, overfitting will be the main issue. Data augmentation is a powerful way to fight overfitting when you're working with image data.
- It's easy to reuse an existing convnet on a new dataset via feature extraction. This is a valuable technique for working with small image datasets.
- As a complement to feature extraction, you can use fine-tuning, which adapts to a new problem some of the representations previously learned by an existing model. This pushes performance a bit further.



Visualizing What ConvNets Learn

- *Visualizing intermediate convnet outputs (intermediate activations)*—Useful for understanding how successive convnet layers transform their input, and for getting a first idea of the meaning of individual convnet filters.
- *Visualizing convnets filters*—Useful for understanding precisely what visual pattern or concept each filter in a convnet is receptive to.
- *Visualizing heatmaps of class activation in an image*—Useful for understanding which parts of an image were identified as belonging to a given class, thus allowing you to localize objects in images.

Reloading Earlier Model

```
>>> model.summary()  <1> As a reminder.
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_5 (Conv2D)	(None, 148, 148, 32)	896
maxpooling2d_5 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_6 (Conv2D)	(None, 72, 72, 64)	18496
maxpooling2d_6 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_7 (Conv2D)	(None, 34, 34, 128)	73856
maxpooling2d_7 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_8 (Conv2D)	(None, 15, 15, 128)	147584
maxpooling2d_8 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_2 (Flatten)	(None, 6272)	0
dropout_1 (Dropout)	(None, 6272)	0
dense_3 (Dense)	(None, 512)	3211776
dense_4 (Dense)	(None, 1)	513

Processing a Single Image

```
img_path = '/Users/fchollet/Downloads/cats_and_dogs_small/test/cats/cat.1700.jpg'  
  
from keras.preprocessing import image  
import numpy as np  
  
img = image.load_img(img_path, target_size=(150, 150))  
img_tensor = image.img_to_array(img)  
img_tensor = np.expand_dims(img_tensor, axis=0)  
img_tensor /= 255.  
  
print(img_tensor.shape)
```

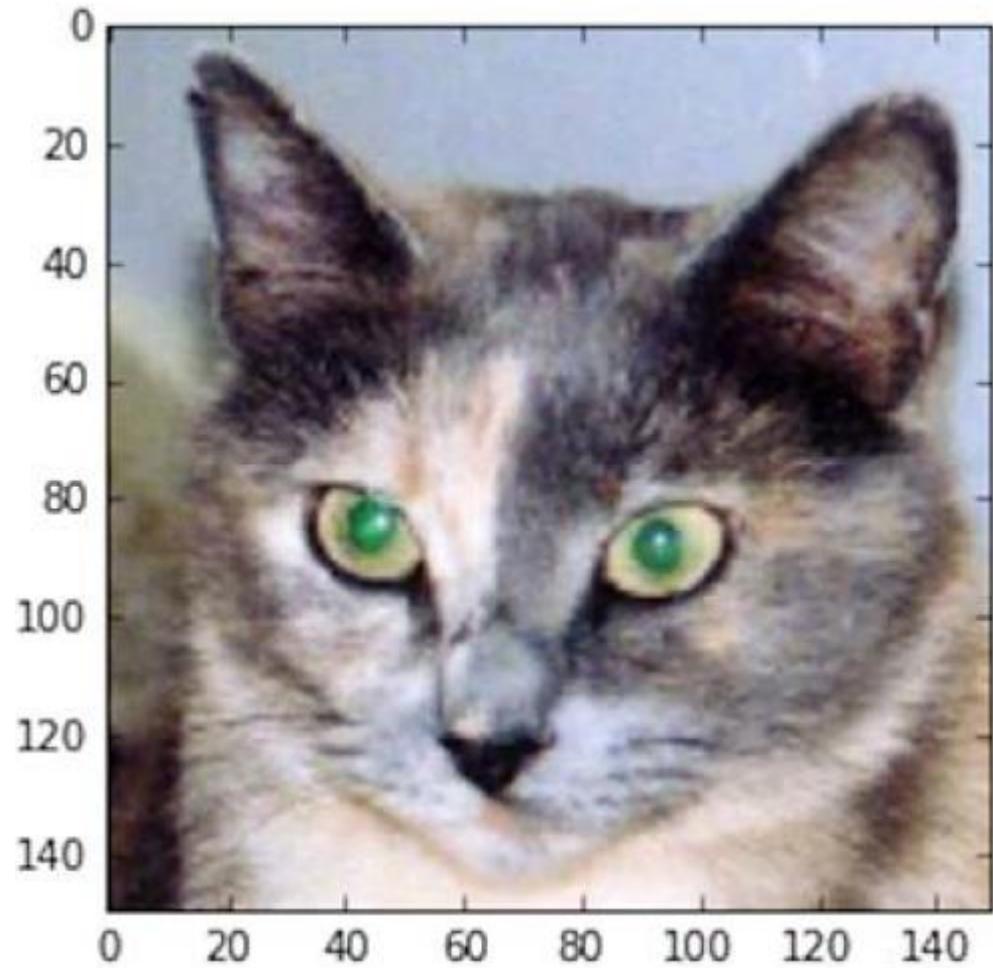
Preprocesses the image into a 4D tensor

Remember that the model was trained on inputs that were preprocessed this way.

Its shape is (1, 150, 150, 3)

Displaying the Image

```
import matplotlib.pyplot as plt  
  
plt.imshow(img_tensor[0])  
plt.show()
```



Creating an Activations Model for the Convolutional Base

```
from keras import models  
→ layer_outputs = [layer.output for layer in model.layers[:8]]  
activation_model = models.Model(inputs=model.input, outputs=layer_outputs) ←
```

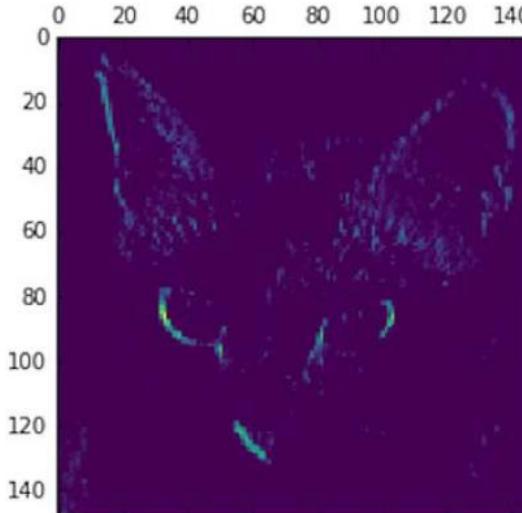
Extracts the outputs of the top eight layers

Creates a model that will return these outputs, given the model input

Visualizing the Activations for the “4th” Convolution Filter of the First Layer

```
activations = activation_model.predict(img_tensor)  
>>> first_layer_activation = activations[0]  
>>> print(first_layer_activation.shape)  
(1, 148, 148, 32)  
  
import matplotlib.pyplot as plt  
  
plt.matshow(first_layer_activation[0, :, :, 4], cmap='viridis')
```

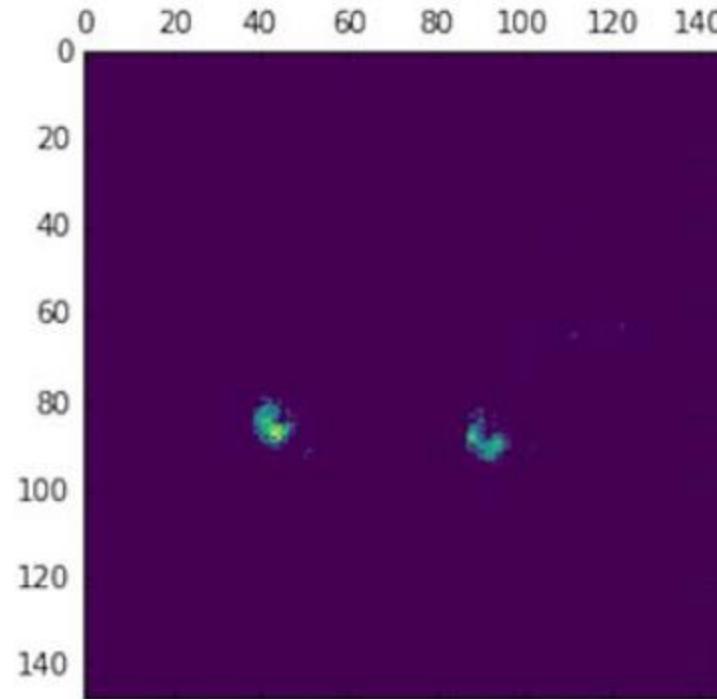
Returns a list of five Numpy arrays: one array per layer activation



Looks like a diagonal edge detector
Check out the nose and ear

Visualizing the Activations for the “7th” Convolution Filter of the First Layer

```
plt.matshow(first_layer_activation[0, :, :, 7], cmap='viridis')
```



Looks like a color detector
Check out the eyes

Visualizing All of the Convolution Activations

```
layer_names = []
for layer in model.layers[:8]:
    layer_names.append(layer.name)

images_per_row = 16

for layer_name, layer_activation in zip(layer_names, activations):
    n_features = layer_activation.shape[-1]
    size = layer_activation.shape[1]
    n_cols = n_features // images_per_row
    display_grid = np.zeros((size * n_cols, images_per_row * size))

    for col in range(n_cols):
        for row in range(images_per_row):
            channel_image = layer_activation[0,
                                              :, :,
                                              col * images_per_row + row]

            channel_image -= channel_image.mean()
            channel_image /= channel_image.std()
            channel_image *= 64
            channel_image += 128
            channel_image = np.clip(channel_image, 0, 255).astype('uint8')
            display_grid[col * size : (col + 1) * size,
                         row * size : (row + 1) * size] = channel_image

    scale = 1. / size
    plt.figure(figsize=(scale * display_grid.shape[1],
                       scale * display_grid.shape[0]))
    plt.title(layer_name)
    plt.grid(False)
    plt.imshow(display_grid, aspect='auto', cmap='viridis')
```

Number of features in the feature map

Tiles the activation channels in this matrix

Post-processes the feature to make it visually palatable

Names of the layers, so you can have them as part of your plot

Displays the feature maps

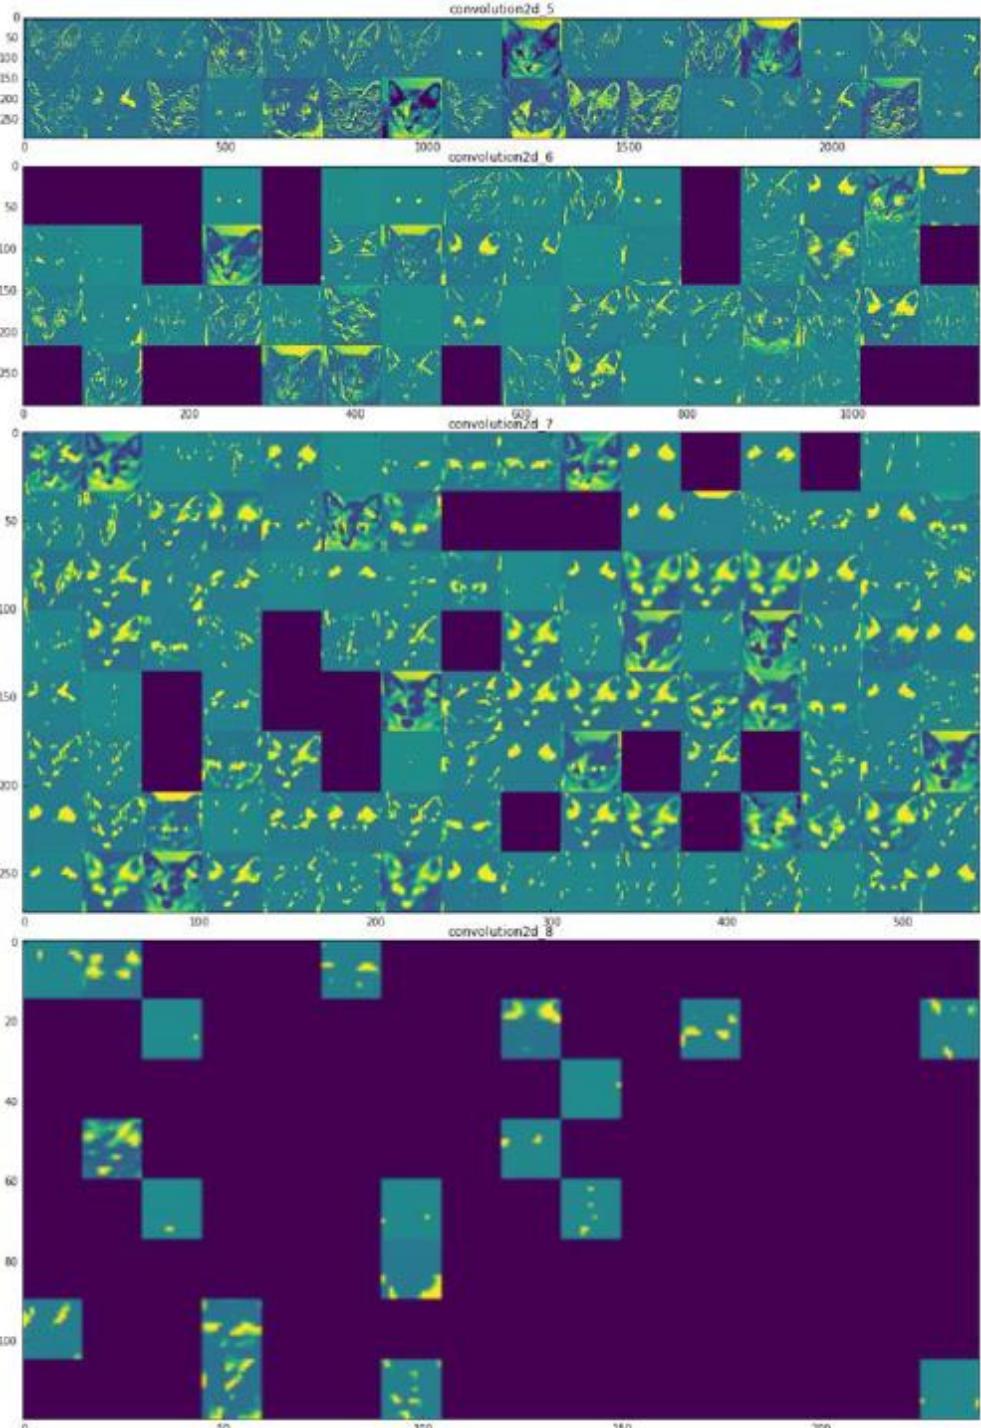
The feature map has shape (l, size, size, n_features).

Tiles each filter into a big horizontal grid

Displays the grid

Activations

- Counts: [32, 64, 128, 128]
- 1st layer acts as a collection of various edge detectors
- as we go from the 1st convolution layer to the 4th convolution layer
 - Features become more abstract (higher-level concepts) and less visually interpretable
 - Activations become more sparse

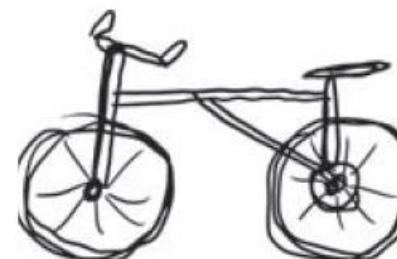


Abstract Representation

- “... a human can remember which abstract objects were present in [a scene] (bicycle, tree) but can’t remember the specific appearance of these objects. In fact, if you tried to draw a generic bicycle from memory, chances are you couldn’t get it even remotely right, even though you’ve seen thousands of bicycles in your lifetime.”



Attempts to draw from memory



Actual picture

Visualizing ConvNet Filters

- We want to visualize an image that maximizes a convolution filter's mean output
- We'll be using gradient ascent on the convolution output
- [better to call it activation rather than loss]

```
from keras.applications import VGG16
from keras import backend as K

model = VGG16(weights='imagenet',
                include_top=False)

layer_name = 'block3_conv1'
filter_index = 0

layer_output = model.get_layer(layer_name).output
loss = K.mean(layer_output[:, :, :, :, filter_index])
```

We're using VGG16

Fetching Gradients

```
grads = K.gradients(loss, model.input)[0] ←  
grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5) ←  
  
iterate = K.function([model.input], [loss, grads])  
  
import numpy as np  
loss_value, grads_value = iterate([np.zeros((1, 150, 150, 3))])
```

The call to gradients returns a list of tensors (of size 1 in this case). Hence, you keep only the first element—which is a tensor.

Add 1e-5 before dividing to avoid accidentally dividing by 0.

Optimizing the Input

Starts from a gray image with some noise

```
→ input_img_data = np.random.random((1, 150, 150, 3)) * 20 + 128.
```

```
step = 1.           ← Magnitude of each gradient update
```

```
for i in range(40):
```

```
→ loss_value, grads_value = iterate([input_img_data])
```

```
input_img_data += grads_value * step
```

Computes the loss value and gradient value

Runs gradient ascent for 40 steps

↑ Adjusts the input image in the direction that maximizes the loss

Gradient ascent to maximize activations:
simply adding the gradients to the image

Converting the Input Tensor to an Image

```
def deprocess_image(x):
    x -= x.mean()
    x /= (x.std() + 1e-5)
    x *= 0.1

    x += 0.5
    x = np.clip(x, 0, 1)

    x *= 255
    x = np.clip(x, 0, 255).astype('uint8')
    return x
```

Normalizes the tensor:
centers on 0, ensures
that std is 0.1

Clips to [0, 1]

Converts to an RGB array

Putting the Pieces Together

Builds a loss function that maximizes the activation of the nth filter of the layer under consideration

```
def generate_pattern(layer_name, filter_index, size=150):
    layer_output = model.get_layer(layer_name).output
    loss = K.mean(layer_output[:, :, :, filter_index])
```

Computes the gradient of the input picture with regard to this loss

Normalization trick: normalizes the gradient

Returns the loss and grads given the input picture

Runs gradient ascent for 40 steps

```
step = 1.
for i in range(40):
    loss_value, grads_value = iterate([input_img_data])
    input_img_data += grads_value * step
```

```
img = input_img_data[0]
return deprocess_image(img)
```

Starts from a gray image with some noise

Visualizing “Optimal” Input for All the Filters

```
layer_name = 'block1_conv1'  
size = 64  
margin = 5  
  
results = np.zeros((8 * size + 7 * margin, 8 * size + 7 * margin, 3))  
  
for i in range(8):  
    for j in range(8):  
        filter_img = generate_pattern(layer_name, i + (j * 8), size=size)  
  
        horizontal_start = i * size + i * margin  
        horizontal_end = horizontal_start + size  
        vertical_start = j * size + j * margin  
        vertical_end = vertical_start + size  
        results[horizontal_start: horizontal_end,  
               vertical_start: vertical_end, :] = filter_img  
  
plt.figure(figsize=(20, 20))  
plt.imshow(results)
```

Empty (black) image to store results

Iterates over the rows of the results grid

Iterates over the columns of the results grid

Generates the pattern for filter $i + (j * 8)$ in `layer_name`

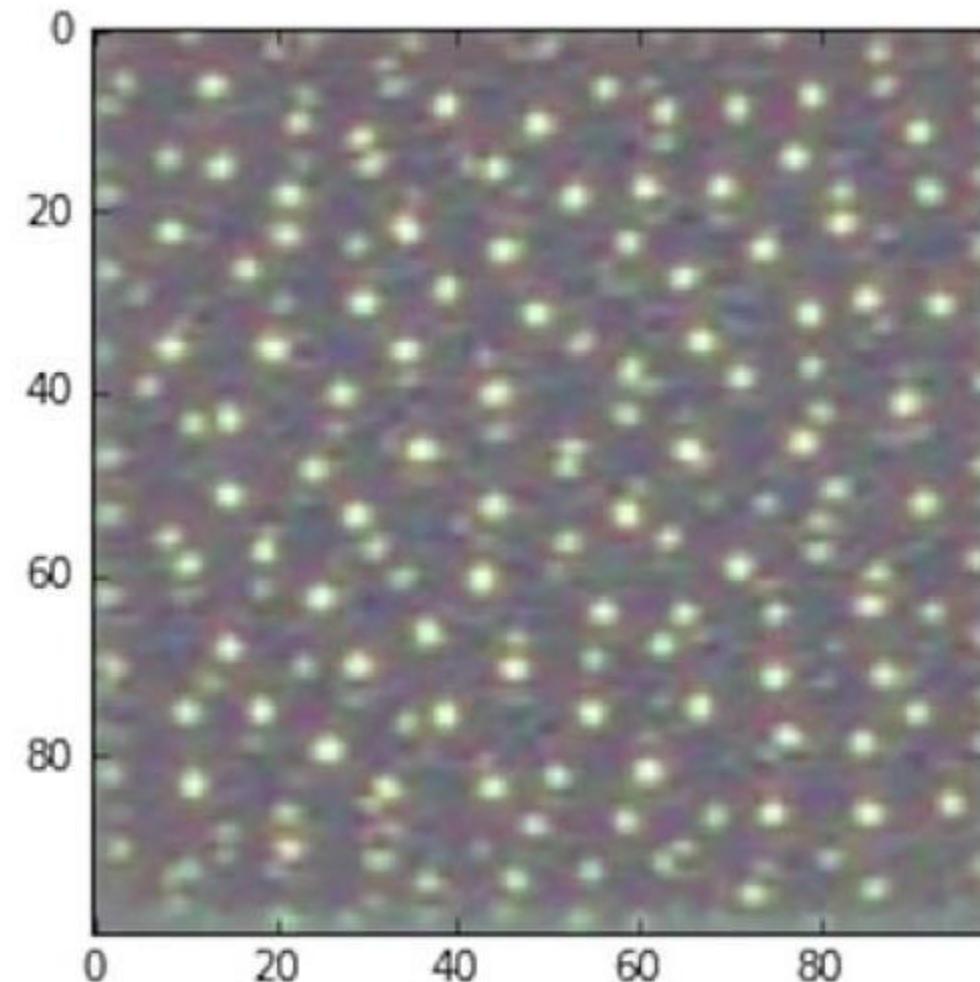
Puts the result in the square (i, j) of the results grid

Displays the results grid

Presenting the Polka Dot Filter ☺

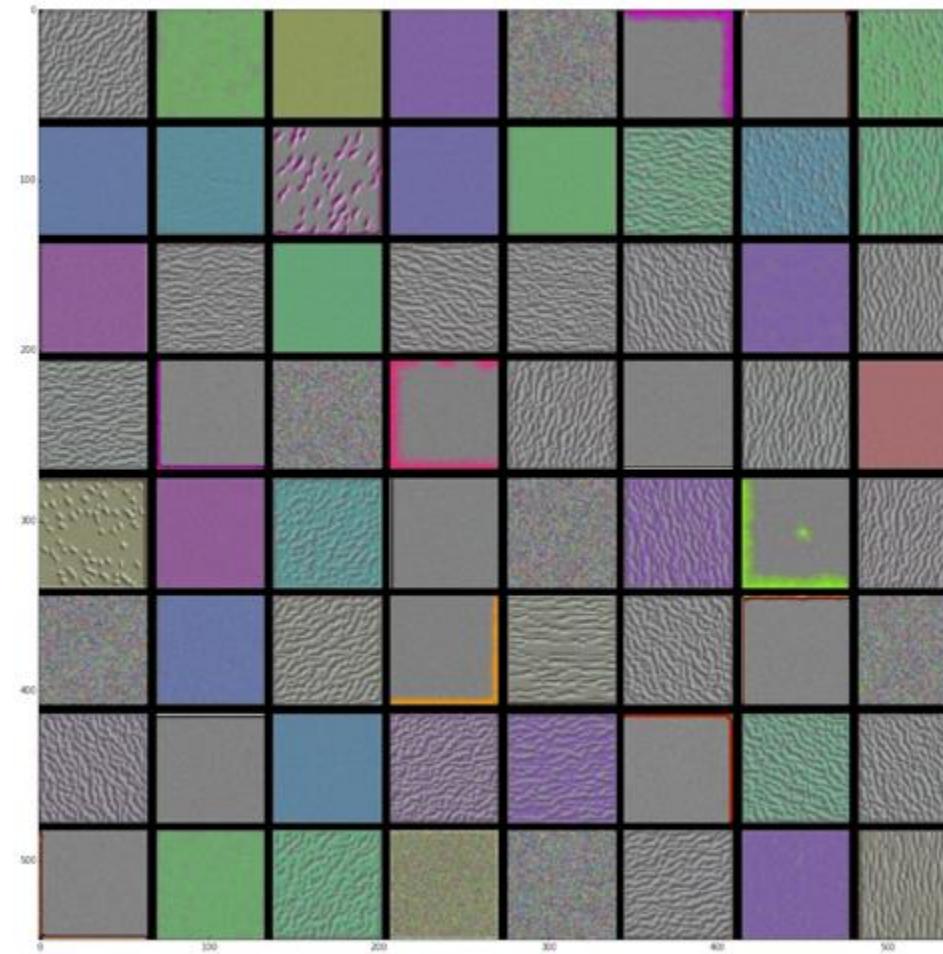
```
>>> plt.imshow(generate_pattern('block3_conv1', 0))
```

from block3_conv1



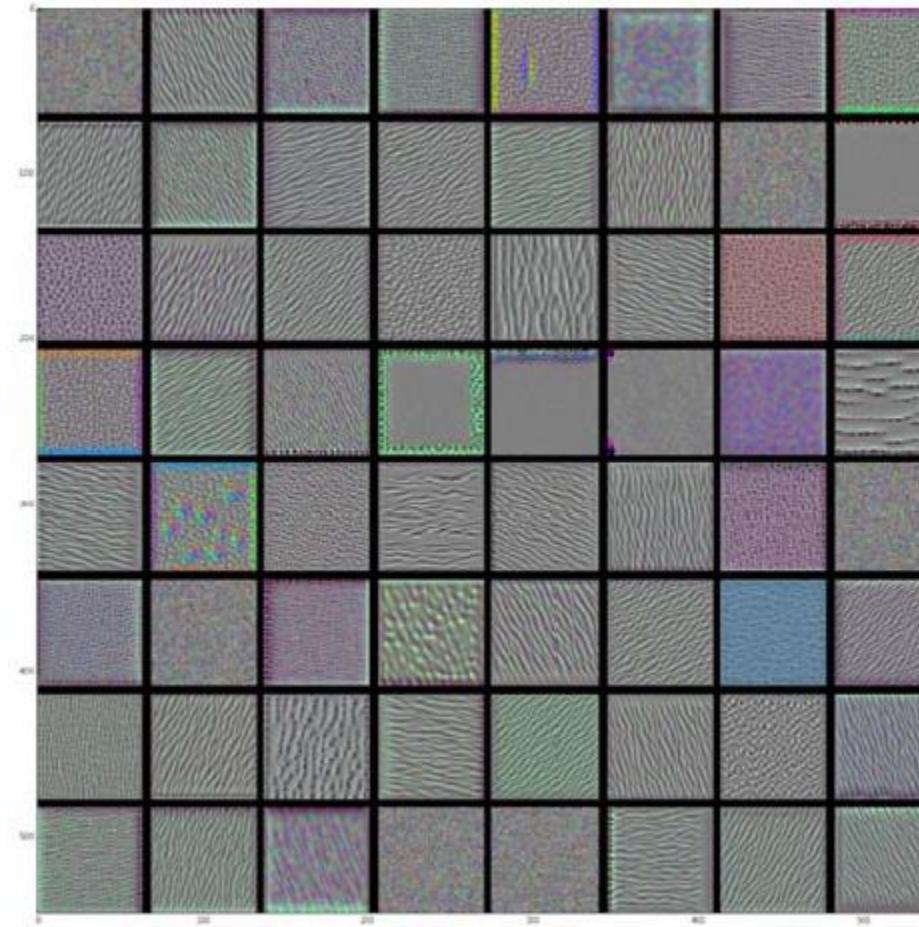
“Optimal” Images for block1_conv1

Simple directional edges and colors



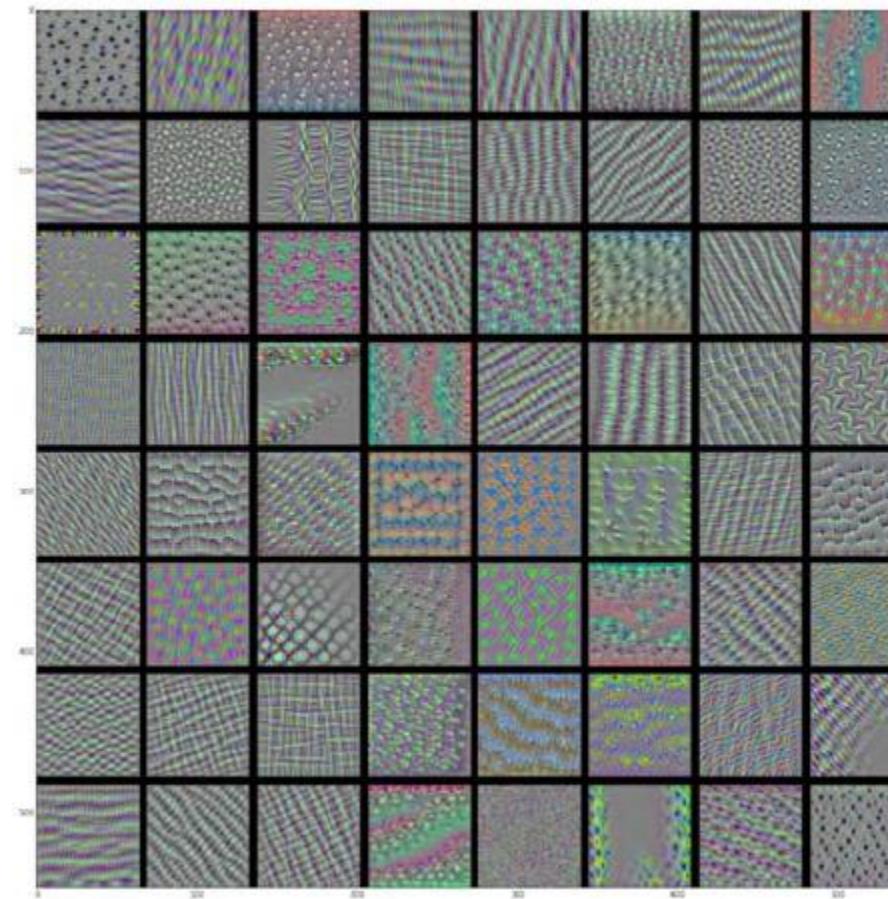
“Optimal” Images for block2_conv1

Simple textures from combinations of edges and colors



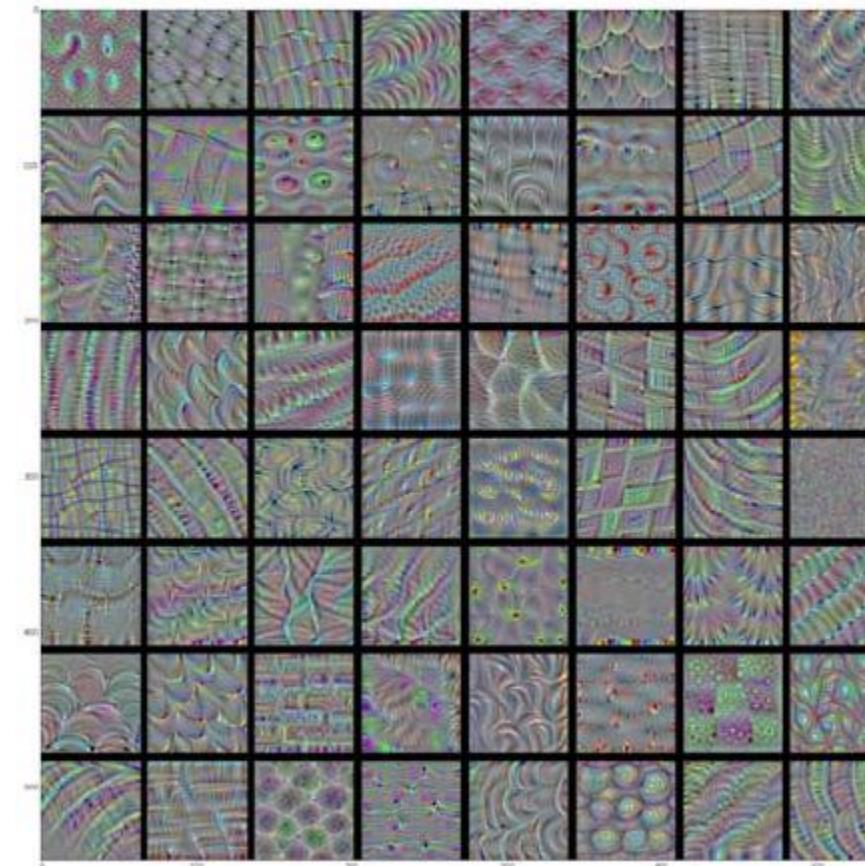
“Optimal” Images for block3_conv1

Includes the polka dot filter



“Optimal” Images for block4_conv1

Higher layer filters are activated by more complex textures, resembling textures found in natural images



Loading the VGG16 Convolutional Base

```
from keras.applications.vgg16 import VGG16  
model = VGG16(weights='imagenet')
```

Note that you include the densely connected classifier on top; in all previous cases, you discarded it.

We've loaded the top as well



Preprocessing an Input Image for VGG16

```
from keras.preprocessing import image  
from keras.applications.vgg16 import preprocess_input, decode_predictions  
import numpy as np
```

```
→ img_path = '/Users/fchollet/Downloads/creativecommons_elephant.jpg'
```

```
→ img = image.load_img(img_path, target_size=(224, 224))
```

```
x = image.img_to_array(img)
```

←
**float32 Numpy array of shape
(224, 224, 3)**

```
x = np.expand_dims(x, axis=0)
```

←
**Adds a dimension to transform the array
into a batch of size (1, 224, 224, 3)**

```
x = preprocess_input(x)
```

←
**Preprocesses the batch (this does
channel-wise color normalization)**

**Python Imaging Library (PIL) image
of size 224 × 224**

Local path to the target image

Using VGG16 to Make Predictions

```
>>> preds = model.predict(x)
>>> print('Predicted:', decode_predictions(preds, top=3) [0])
Predicted: [(u'n02504458', u'African_elephant', 0.92546833),
(u'n01871265', u'tusker', 0.070257246),
(u'n02504013', u'Indian_elephant', 0.0042589349)]
```

The top three classes predicted for this image are as follows:

- African elephant (with 92.5% probability)
- Tusker (with 7% probability)
- Indian elephant (with 0.4% probability)

```
>>> np.argmax(preds [0])
```

386

Gradients of the Elephant Softmax Activation with respect to the block5_conv3 Activation

“African elephant” entry in the prediction vector

```
→ african_elephant_output = model.output[:, 386]  
last_conv_layer = model.get_layer('block5_conv3')
```

Output feature map of the block5_conv3 layer, the last convolutional layer in VGG16

Gradient of the “African elephant” class with regard to the output feature map of block5_conv3

Vector of shape (512,), where each entry is the mean intensity of the gradient over a specific feature-map channel

```
→ grads = K.gradients(african_elephant_output, last_conv_layer.output)[0]  
pooled_grads = K.mean(grads, axis=(0, 1, 2))
```

Generating the Heatmap for block5_conv3 Activations

```
iterate = K.function([model.input],  
                    [pooled_grads, last_conv_layer.output[0]])  
  
→ pooled_grads_value, conv_layer_output_value = iterate([x])  
  
for i in range(512):  
    conv_layer_output_value[:, :, i] *= pooled_grads_value[i]  
  
heatmap = np.mean(conv_layer_output_value, axis=-1) ←
```

Values of these two quantities, as Numpy arrays, given the sample image of two elephants

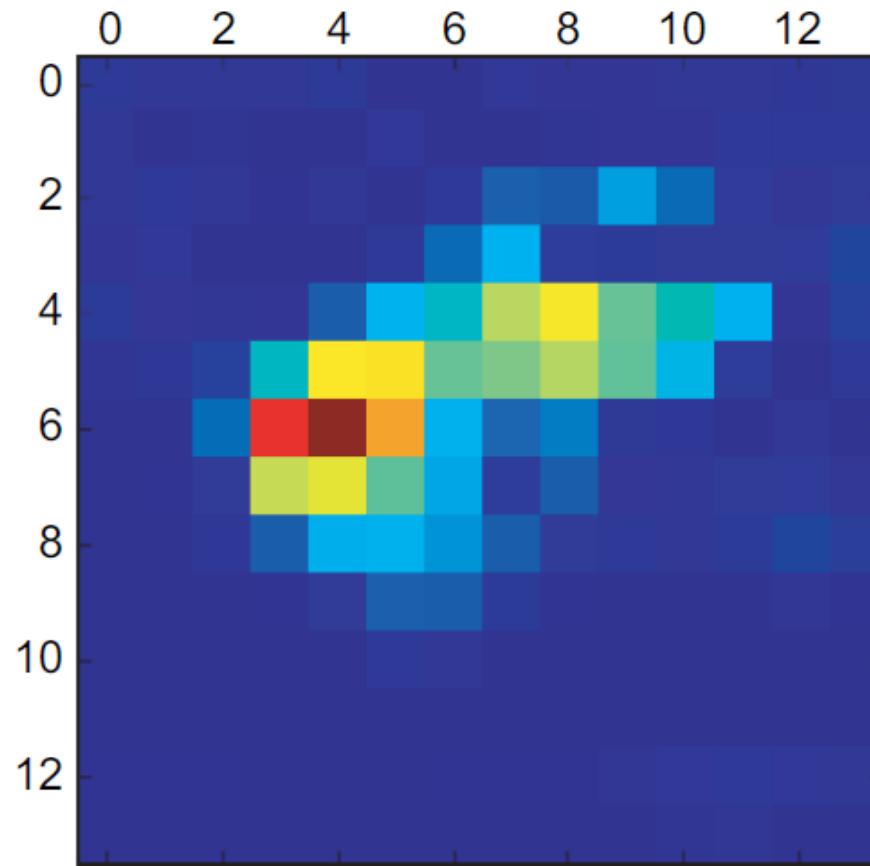
Lets you access the values of the quantities you just defined: pooled_grads and the output feature map of block5_conv3, given a sample image

The channel-wise mean of the resulting feature map is the heatmap of the class activation.

Multiplies each channel in the feature-map array by “how important this channel is” with regard to the “elephant” class

Normalizing and Displaying the Heatmap

```
heatmap = np.maximum(heatmap, 0)
heatmap /= np.max(heatmap)
plt.matshow(heatmap)
```



Superimposing the Heatmap onto the Original Picture

```
import cv2
img = cv2.imread(img_path)           ← Uses cv2 to load the original image
heatmap = cv2.resize(heatmap, (img.shape[1], img.shape[0])) ← Resizes the heatmap to be the same size as the original image
heatmap = np.uint8(255 * heatmap)      ← Converts the heatmap to RGB
heatmaps = cv2.applyColorMap(heatmap, cv2.COLORMAP_JET)
superimposed_img = heatmap * 0.4 + img
cv2.imwrite('/Users/fchollet/Downloads/elephant_cam.jpg', superimposed_img) ← Saves the image to disk
```

0.4 here is a heatmap intensity factor.

Applies the heatmap to the original image

Class Activation Heatmap Superimposed



This visualization technique answers two important questions:

- Why did the network think this image contained an African elephant?
- Where is the African elephant located in the picture?

Chapter Summary

- Convnets are the best tool for attacking visual-classification problems.
- Convnets work by learning a hierarchy of modular patterns and concepts to represent the visual world.
- The representations they learn are easy to inspect—convnets are the opposite of black boxes!
- You’re now capable of training your own convnet from scratch to solve an image-classification problem.
- You understand how to use visual data augmentation to fight overfitting.
- You know how to use a pretrained convnet to do feature extraction and fine-tuning.
- You can generate visualizations of the filters learned by your convnets, as well as heatmaps of class activity.

Runtime Complexity

- When describing the runtime complexity of a computation, we're allowed to ignore constants; so if a computation involves $2*k+1$ operations, we would say the runtime complexity is $O(k)$
 - Big 'O' ["oh"]: an upper-bound on the number of operations
 - Big ' Ω ' ["omega"]: a lower-bound on the number of operations
 - Big ' Θ ' ["theta"]: both a lower-bound and an upper-bound on the number of operations

Runtime Complexity for a Dense Layer

- Reminder: a Dense layer neuron has a weight for every input feature from the previous layer, so Dense layer neurons are said to generate “global” features
- Runtime (inference) complexity for a dense layer is the same as runtime complexity for matrix multiplication:
 $\text{num_input_observations} * \text{num_output_features} * \text{num_input_features}$
- So for MNIST training images with a dense layer of 512 neurons, the runtime complexity is $60,000 * 512 * 784 = 24,084,480,000$
we’re essentially ignoring the addition operations in our count, as the additional time for addition is somewhat negligible

Runtime Complexity for a Convolutional Layer

- Reminder: a Conv2D layer neuron has weights for only a small neighborhood of input pixels, so Conv2D layer neurons are said to generate “local” features
- Runtime complexity for a Conv2D layer involves terms for the number of neighborhoods evaluated as well as the size of the neighborhood:

Same filter can detect pattern in a different image location

ImageCount

* FilterCount

$$\left(\left(\frac{\text{InputImageHeight} - \text{FilterHeight} + 2 * \text{Padding}}{\text{StrideHeight}} + 1 \right) * \left(\frac{\text{InputImageWidth} - \text{FilterWidth} + 2 * \text{Padding}}{\text{StrideWidth}} + 1 \right) \right)$$

* ($\text{FilterHeight} * \text{FilterWidth} * \text{InputChannels}$)

- So for MNIST training images with a Conv2D layer of 64 3x3 filters, 1x1 stride, and “same” padding [padding = (filter_size – 1)/2 = 1], the runtime complexity is:
 $60,000 * 64 * 28 * 28 * (3 * 3 * 1) = 27,095,040,000$
- What happens to the runtime complexity if there are 64 input channels?

Notebooks and Other Examples

- Deep Learning with Python

<https://github.com/fchollet/deep-learning-with-python-notebooks>

- Deep Learning Illustrated

<https://github.com/the-deep-learners/deep-learning-illustrated>

- <https://keras.io/examples/>

PyTorch

- PyTorch, the other major neural network implementation framework, refers to a Dense layer as a Linear layer:
<https://discuss.pytorch.org/t/pytorch-equivalent-of-keras/29412>
- While the implementation details vary somewhat, the underlying concepts for Tensorflow and PyTorch are the same
“shields up” ☺